

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT **ON** **MACHINE LEARNING**

Submitted by:

MAHAVIR NAHATA (1BM21CS100)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU-560019

March 2024-June 2024



B. M. S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering

CERTIFICATE

This is to certify that the Lab work entitled “**MACHINE LEARNING**” carried out by **MAHAVIR NAHATA(1BM21CS100)**, who is bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2023-24. The Lab report has been approved as it satisfies the academic requirements in respect of Machine Learning Lab - **(22CS3PCMAL)** work prescribed for the said degree.

Sunayana S

Assistant Professor
Department of CSE
BMSCE, Bengaluru

Dr. Jyothi S Nayak

Professor and Head
Department of CSE
BMSCE, Bengaluru

INDEX

Sl. No.	Experiment Title	Page No.
1	Write a python program to import and export data using Pandas library functions	3-5
2	Demonstrate various data pre-processing techniques for a given dataset	6-11
3	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	12-18
4	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	19-25
5	Build Logistic Regression Model for a given dataset	26-30
6	Build KNN Classification model for a given dataset.	31-34
7	Build Support vector machine model for a given dataset	35-38
8	Build Artificial Neural Network model with back propagation on a given dataset	39-41
9	a)Implement Random forest ensemble method on a given dataset. b)Implement Boosting ensemble method on a given dataset.	42-51
10	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	52-55
11	Implement Dimensionality reduction using Principle Component Analysis (PCA) method.	56-59

PROGRAM 1

Date:05-04-2024

Write a python program to import and export data using Pandas library functions

05/04/24

Program -1

Write a python program to import and export data using Pandas library functions

IMPORT:

```
import pandas as pd
airbnb_data = pd.read_csv("listings.csv")
airbnb_data.head()
```

EXPORT:

```
airbnb_data.to_csv("exported-listings.csv")
```

READING DATA FROM URL:

```
url = "https://archive.ics.uci.edu/ml/machine-learning-  
-databases/iris/iris.data"  
col_names = ["sepal-length-in-cm", "sepal-width-  
in-cm", "petal-length-in-cm",  
"petal-width-in-cm", "class"]  
iris_data = pd.read_csv(url, names=col_names)  
iris_data.head()
```

Import:

import pandas as pd

Read the CSV file

airbnb_data = pd.read_csv("listings.csv")

View the first 5 rows

airbnb_data.head()

	id	name	host_id	host_name	neighbourhood_group	neighbourhood	latitude	longitude	room_type	price	minimum_nights	number_of_reviews	last_n
0	329172	Hillside designer home, 10 min. drive to	1680871	Janet	NaN	78746	30.30085	-97.80794	Entire home/apt.	495	3	7	2022-
1	329306	Urban Homestead, 5 minutes to downtown	880571	Angel	NaN	78702	30.27232	-97.72579	Private room	63	2	570	2022-
2	331549	One Room with Private Bathroom	1690383	Sandra	NaN	78725	30.23911	-97.58625	Private room	100	2	0	
3	333815	Solar Sanctuary - Austin Room	372962	Kim	NaN	78704	30.25381	-97.75262	Private room	102	2	164	2022-
4	333442	Rare Secluded 1940s Estate	1698318	Virginia	NaN	78703	30.31267	-97.76641	Entire home/apt.	286	3	163	2022-

Export:

airbnb_data.to_csv("exported_listings.csv")

	id	name	host_id	host_name	neighbourhood_group	neighbourhood	latitude	longitude	room_type	price	minimum_nights	number_of_reviews	last_n
0	329172	Hillside designer home, 10 min. drive to	1680871	Janet	NaN	78746	30.30085	-97.80794	Entire home/apt.	495	3	7	2022-
1	329306	Urban Homestead, 5 minutes to downtown	880571	Angel	NaN	78702	30.27232	-97.72579	Private room	63	2	570	2022-
2	331549	One Room with Private Bathroom	1690383	Sandra	NaN	78725	30.23911	-97.58625	Private room	100	2	0	
3	333815	Solar Sanctuary - Austin Room	372962	Kim	NaN	78704	30.25381	-97.75262	Private room	102	2	164	2022-
4	333442	Rare Secluded 1940s Estate	1698318	Virginia	NaN	78703	30.31267	-97.76641	Entire home/apt.	286	3	163	2022-

Reading data from URL:

```
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
```

```
# Define the column names
```

```
col_names = ["sepal_length_in_cm",  
             "sepal_width_in_cm",  
             "petal_length_in_cm",  
             "petal_width_in_cm",  
             "class"]
```

```
# Read data from URL
```

```
iris_data = pd.read_csv(url, names=col_names)
```

```
iris_data.head()
```

```
[10]:
```

	sepal_length_in_cm	sepal_width_in_cm	petal_length_in_cm	petal_width_in_cm	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

PROGRAM 2

Date:05-04-2024

Demonstrate various data pre-processing techniques for a given dataset

Code and Output

2. Importing and Exploration of the dataset

```
In [2]: # Loading the data and setting the unique client_id as the index::  
df = pd.read_csv('/content/loans.csv', index_col = 'client_id')
```

```
In [3]: # # showing the first 5 rows of the dataset:  
df.head()
```

```
Out[3]:
```

	loan_type	loan_amount	repaid	loan_id	loan_start	loan_end	rate
client_id							
46109	home	13672	0	10243	2002-04-16	2003-12-20	2.15
46109	credit	9794	0	10984	2003-10-21	2005-07-17	1.25
46109	home	12734	1	10990	2006-02-01	2007-07-05	0.68
46109	cash	12518	1	10596	2010-12-08	2013-05-05	1.24
46109	credit	14049	1	11415	2010-07-07	2012-05-21	3.13

```
In [4]: # To check the Dimensions of the dataset:  
df.shape
```

```
Out[4]: (443, 7)
```

```
In [5]: # Checking the Info of the data:  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 443 entries, 46109 to 26945  
Data columns (total 7 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   loan_type    443 non-null    object  
1   loan_amount  443 non-null    int64  
2   repaid       443 non-null    int64  
3   loan_id      443 non-null    int64  
4   loan_start   443 non-null    object  
5   loan_end     443 non-null    object  
6   rate         443 non-null    float64  
dtypes: float64(1), int64(3), object(3)  
memory usage: 27.7+ KB
```

3. Checking the datatypes of the columns

```
In [6]: df.dtypes
```

```
Out[6]: loan_type      object
loan_amount    int64
repaid         int64
loan_id        int64
loan_start     object
loan_end       object
rate          float64
dtype: object
```

4. Converting the data types of columns

- loan_id to object
- repaid to category dtype
- loan_start and loan_end to date type

```
In [7]: # loan_id:
df['loan_id'] = df['loan_id'].astype('object')

# repaid:
df['repaid'] = df['repaid'].astype('category')
```

```
In [8]: # loan_start:
df['loan_start'] = pd.to_datetime(df['loan_start'], format = '%Y-%m-%d')

# loan_end:
df['loan_end'] = pd.to_datetime(df['loan_end'], format = '%Y-%m-%d')
```

Checking the datatypes again:

```
In [9]: df.dtypes
```

```
Out[9]: loan_type      object
loan_amount    int64
repaid         category
loan_id        object
loan_start     datetime64[ns]
loan_end       datetime64[ns]
rate          float64
dtype: object
```

5. Summary Statistics of the data

```
In [10]: # Summary Statistics for Numerical data:
df.describe()
```

```
Out[10]:
```

	loan_amount	loan_start	loan_end	rate
count	443.000000	443	443	443.000000
mean	7982.311312	2007-08-02 12:56:53.092550912	2009-08-23 11:29:37.246049536	3.217156
min	559.000000	2000-01-26 00:00:00	2001-08-02 00:00:00	0.010000
25%	4232.500000	2003-10-19 00:00:00	2005-09-12 12:00:00	1.220000
50%	8320.000000	2007-03-10 00:00:00	2008-03-19 00:00:00	2.780000
75%	11739.000000	2011-07-31 00:00:00	2013-09-11 12:00:00	4.750000
max	14971.000000	2014-11-11 00:00:00	2017-05-07 00:00:00	12.620000
std	4172.891992	NaN	NaN	2.397168


```
In [11]: # Summary Statistics for Categorical data:
df.describe(exclude=[np.number])
```

```
Out[11]:
```

	loan_type	repaid	loan_id	loan_start	loan_end
count	443	443.0	443.0	443	443
unique	4	2.0	443.0	NaN	NaN
top	home	1.0	10243.0	NaN	NaN
freq	121	237.0	1.0	NaN	NaN
mean	NaN	NaN	NaN	2007-08-02 12:56:53.092550912	2009-08-23 11:35:37.246049536
min	NaN	NaN	NaN	2000-01-26 00:00:00	2001-08-02 00:00:00
25%	NaN	NaN	NaN	2003-10-19 00:00:00	2005-09-12 12:00:00
50%	NaN	NaN	NaN	2007-03-10 00:00:00	2009-03-19 00:00:00
75%	NaN	NaN	NaN	2011-07-31 00:00:00	2013-09-11 12:00:00
max	NaN	NaN	NaN	2014-11-11 00:00:00	2017-05-07 00:00:00

6. Missing Values

```
In [12]: # use isnull().sum() to check for missing values
df.isnull().sum()
```

```
Out[12]: loan_type      0
loan_amount    0
repaid         0
loan_id        0
loan_start     0
loan_end       0
rate           0
dtype: int64
```


There are no missing values in the data.

Sk-learn library has an in-built function called Iterative Imputer to impute the missing values. Its sklearn documentation: <https://scikit-learn.org/stable/modules/generated/sklearn.impute.IterativeImputer.html>


7. Outliers Treatment

To check for the presence of outliers, we plot Boxplot.

```
In [13]: # For loan_amount
df['loan_amount'].plot(kind='box')
plt.show()
```



```
In [18]: # for rate
df['rate'].plot(kind='box')
plt.show()
```



We can see that there are no outliers in the loan_amount column and some outliers are present in the rate column. To treat for outliers can either cap the values or transform the data. Shall demonstrate both the approaches here.

8. Transformation

8a. SQRT transformation

```
In [19]: df['SQRT_RATE'] = df['rate']**0.5
```

```
In [20]: df['sqrt_rate'] = np.sqrt(df['rate'])
```

```
In [21]: df.head()
```

```
Out[21]:
```

	loan_type	loan_amount	repaid	loan_id	loan_start	loan_end	rate	SQRT_RATE	sqrt_rate
client_id									
46109	home	13672	0	10043	2000-04-16	2005-12-20	2.15	1.466288	1.466288
46109	credit	9794	0	10984	2000-10-21	2005-07-17	1.25	1.118034	1.118034
46109	home	12734	1	10990	2006-02-01	2007-07-08	0.88	0.938281	0.938281
46109	cash	12518	1	10996	2010-12-06	2013-05-08	1.24	1.113553	1.113553
46109	credit	14049	1	11415	2010-07-07	2012-05-21	3.13	1.769181	1.769181

```
In [44]: #checking the skewness, kurtosis between the original and transformed data:
print("The skewness of the original data is {}".format(df.rate.skew()))
print("The skewness of the SQRT transformed data is {}".format(df.sqRT_RATE.skew()))

print("")

print("The kurtosis of the original data is {}".format(df.rate.kurt()))
print("The kurtosis of the SQRT transformed data is {}".format(df.sqRT_RATE.kurt()))
```


The skewness of the original data is 0.354294634329949
The skewness of the SQRT transformed data is 0.04664354695520862

The kurtosis of the original data is 0.42437165343736433
The kurtosis of the SQRT transformed data is -0.6315437642982899

```
In [45]: # plotting the distribution

fig, axes = plt.subplots(1,2, figsize=(18,5))
sns.distplot(df['rate'], ax=axes[0])
sns.distplot(df['sqRT_RATE'], ax=axes[1])

plt.show()
```



Result:

The Rate column was right skewed earlier. The skewness and kurtosis as reduced significantly. The transformed SQRT rate, on the right graph resembles normal distribution now.

8b. Log Transformation

```
In [46]: df['Log_Rate'] = np.log(df['rate'])
```

```
In [47]: df.head()
```

```
Out[47]:
```

	loan_type	loan_amount	repaid	loan_id	loan_start	loan_end	rate	SQRT_RATE	sqRT_RATE	Log_Rate
client_id										
46109	home	10672	0	10243	2003-04-16	2003-12-22	2.15	1.466288	1.466288	0.763465
46109	credit	8794	0	10864	2003-10-21	2006-07-17	1.25	1.118034	1.118034	0.223144
46109	home	12734	1	10890	2004-03-01	2007-07-08	0.85	0.921921	0.921921	-0.283662
46109	car	12516	1	10646	2010-12-06	2013-05-05	1.24	1.113553	1.113553	0.219111
46109	credit	14249	1	11415	2010-07-07	2012-05-21	5.15	2.269161	2.269161	1.341035

```
In [48]: print("The skewness of the original data is {}".format(df.rate.skew()))
print("The skewness of the SQRT transformed data is {}".format(df.sqRT_RATE.skew()))
print("The skewness of the LOG transformed data is {}".format(df['Log_Rate'].skew()))

print("")

print("The kurtosis of the original data is {}".format(df.rate.kurt()))
print("The kurtosis of the SQRT transformed data is {}".format(df.sqRT_RATE.kurt()))
print("The kurtosis of the LOG transformed data is {}".format(df['Log_Rate'].kurt()))
```

The skewness of the original data is 0.354294634329949
The skewness of the SQRT transformed data is 0.04664354695520862
The skewness of the LOG transformed data is -1.594521763232392


The kurtosis of the original data is 0.42437165343736433
The kurtosis of the SQRT transformed data is -0.6315437642982899
The kurtosis of the LOG transformed data is 4.327632581953235

```
In [11]: # plot the graph:

fig, axes = plt.subplots(1,3,figsize=(15,5))

sns.distplot(df['rate'], ax=axes[0])
sns.distplot(df['sqrt_rate'], ax=axes[1])
sns.distplot(df['log_rate'], ax=axes[2])

plt.show()
```



Inference:

Log Transformation made the rate left skewed and more peaked.

However, Log transformation is more closer to 0 and hence is more normal. Though it heavily manipulates the data.

In our case, square root transformation is more suitable.

```
In [14]: # using lambda function /

df['LOG_Rate'] = df['rate'].apply(lambda x:np.log(x))
```

```
In [15]: df.head()
```

```
Out[15]:
```

	loan_type	loan_amount	repaid	loan_id	loan_start	loan_end	rate	SQRT_RATE	sqrt_rate	Log_Rate	LOG_Rate
client_id											
46109	home	13672	0	10243	2002-04-16	2003-12-20	2.15	1.466288	1.466288	0.765468	0.765468
46109	credit	9794	0	10994	2003-10-21	2005-07-17	1.25	1.118034	1.118034	0.223144	0.223144
46109	home	12734	1	10990	2006-02-01	2007-07-05	0.68	0.824621	0.824621	-0.389662	-0.389662
46109	cash	12518	1	10996	2010-12-08	2013-08-05	1.24	1.113553	1.113553	0.215111	0.215111
46109	credit	14049	1	11415	2010-07-07	2012-05-21	3.13	1.769181	1.769181	1.141033	1.141033

PROGRAM 3

Date: 12-04-2024

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample

Algorithm:

12/04/24

Program - 2
Decision Tree ID3

Algorithm:

ID3 (Examples, Target-attribute, Attributes)

- Create a Root node for the tree
- Return If all Examples are positive, Return the single-node tree Root, with Label = +
- If all Examples are negative, Return the single-node tree Root, with Label = -
- If attribute u empty, Return the single-node tree Root, with Label = most common value of Target-attribute in Examples
- Otherwise Begin:
 - $A \leftarrow$ the attribute from Attributes that best classifies Examples.
 - The decision attribute for Root $\leftarrow A$.
 - For each possible value, v_i , of A,
 - Add a new tree branch below Root, corresponding to the test $A = v_i$
 - Let Examples $_{v_i}$ be the subset of Examples that have values v_i for A.
 - If Examples $_{v_i}$ is empty.
 - Then below this new branch add a leaf node with Label = most common value of Target-attribute
 - Else below this new branch add the subtree ID3.
- End
- Return Root

Code:

Importing Database

```
In [1]: # Importing the required libraries
import pandas as pd
import numpy as np
import math

# Reading the dataset (Tennis-dataset)
data = pd.read_csv('/content/PlayTennis.csv')
```

```
In [2]: from google.colab import drive
drive.mount('/content/drive')
```

```
In [3]: def highlight(cell_value):
    """
    Highlight yes / no values in the dataframe
    """
    color_1 = 'background-color: pink;'
    color_2 = 'background-color: lightgreen;'

    if cell_value == 'no':
        return color_1
    elif cell_value == 'yes':
        return color_2

data.style.applymap(highlight)\
.set_properties(subset=data.columns, **{'width': '100px'})\
.set_table_styles([{'selector': 'th', 'props': [('background-color', 'lightgray'), ('border', '1px solid gray'), ('font-weight', 'bold')]},
                    {'selector': 'tr:hover', 'props': [('background-color', 'white'), ('border', '1.5px solid black')]}])
```

```
Out[1]:
```

	outlook	temp	humidity	windy	play
0	sunny	hot	high	False	no
1	sunny	hot	high	True	no
2	overcast	hot	high	False	yes
3	rainy	mild	high	False	yes
4	rainy	cool	normal	False	yes
5	rainy	cool	normal	True	no
6	overcast	cool	normal	True	yes
7	sunny	mild	high	False	no
8	sunny	cool	normal	False	yes
9	rainy	mild	normal	False	yes
10	sunny	mild	normal	True	yes
11	overcast	mild	high	True	yes
12	overcast	hot	normal	False	yes
13	rainy	mild	high	True	no

Entropy of the dataset

```
In [4]: def find_entropy(data):
    """
    Returns the entropy of the class or features
    formula:  $- \sum P(X) \log P(X)$ 
    """
    entropy = 0
    for i in range(data.unique()):
        x = data.value_counts()[i]/data.shape[0]
        entropy += (- x * math.log(x,2))
    return round(entropy,3)

def information_gain(data, data_):
    """
    Returns the information gain of the features
    """
    info = 0
    for i in range(data_.unique()):
        df = data[data_ == data_.unique()[i]]
        w_avg = df.shape[0]/data.shape[0]
        entropy = find_entropy(df.play)
        x = w_avg * entropy
        info += x
    ig = find_entropy(data.play) - info
    return round(ig, 3)

def entropy_and_infogain(datas, feature):
    """
    Grouping features with the same class and computing their
    entropy and information gain for splitting
    """
    for i in range(data[feature].unique()):
        df = datas[datas[feature]==data[feature].unique()[i]]
        if df.shape[0] < 1:
            continue

        display(df[[feature, 'play']].style.applymap(highlight)\
                .set_properties(subset=[feature, 'play'], **{'width': '80px'})\
                .set_table_styles([{'selector': 'th', 'props': [('background-color', 'lightgray'),\
                                                                ('border', '1px solid gray'),\
                                                                ('font-weight', 'bold')]},\
                                   {'selector': 'td', 'props': [('border', '1px solid gray')]},\
                                   {'selector': 'tr:hover', 'props': [('background-color', 'white'),\
                                                                ('border', '1.5px solid black')]}]))

        print(f'Entropy of {feature} - {data[feature].unique()[i]} = {find_entropy(df.play)}')
        print(f'Information Gain for {feature} = {information_gain(datas, datas[feature])}')

In [5]: print(f'Entropy of the entire dataset: {find_entropy(data.play)}')
```

Entropy of the entire dataset: 0.94

Entropy and Information Gain of temperature

```
In [4]: entropy_and_infogain(data, 'temp')
```

	temp	play
0	hot	no
1	hot	no
2	hot	yes
12	hot	yes

Entropy of temp - hot = 1.0

	temp	play
3	mild	yes
7	mild	no
9	mild	yes
10	mild	yes
11	mild	yes
13	mild	no

Entropy of temp - mild = 0.918

	temp	play
4	cool	yes
5	cool	no
6	cool	yes
8	cool	yes

Entropy of temp - cool = 0.811
Information Gain for temp = 0.829

Entropy and Information Gain of humidity

```
In [7]: entropy_and_infogain(data, 'humidity')
```

	humidity	play
0	high	no
1	high	no
2	high	yes
3	high	yes
7	high	no
11	high	yes
13	high	no

Entropy of humidity - high = 0.985

	humidity	play
4	normal	yes
5	normal	no
6	normal	yes
8	normal	yes
9	normal	yes
10	normal	yes
12	normal	yes

Entropy of humidity - normal = 0.592
Information Gain for humidity = 0.151

Entropy and Information Gain of windy

```
In [8]: entropy_and_infogain(data, "windy")
```

	windy	play
--	-------	------

0	False	no
---	-------	----

2	False	yes
---	-------	-----

3	False	yes
---	-------	-----

4	False	yes
---	-------	-----

7	False	no
---	-------	----

8	False	yes
---	-------	-----

9	False	yes
---	-------	-----

12	False	yes
----	-------	-----

Entropy of windy - False = 0.811

	windy	play
--	-------	------

1	True	no
---	------	----

5	True	no
---	------	----

6	True	yes
---	------	-----

10	True	yes
----	------	-----

11	True	yes
----	------	-----

13	True	no
----	------	----

Entropy of windy - True = 1.0

Information Gain for windy = 0.048

Rainy Outlook

Rainy -outlook

```
In [9]: rainy = data[data['outlook'] == "rainy"]
rainy.style.applymap(highlight)
rainy.set_properties(subset=data.columns, **{"width": "100px"})
rainy.set_table_styles([{"selector": "th", "props": [{"background-color": "lightgray"}, {"border": "1px solid gray"}, {"font-weight": "bold"}]}, {"selector": "tr:hover", "props": [{"background-color": "white"}, {"border": "1.5px solid black"}]})
```

```
Out[9]:
```

	outlook	temp	humidity	windy	play
3	rainy	mild	high	False	yes
4	rainy	cool	normal	False	yes
5	rainy	cool	normal	True	no
9	rainy	mild	normal	False	yes
13	rainy	mild	high	True	no

```
In [10]: print(f'Entropy of the Rainy dataset: {find_entropy(rainy.play)}')
```

Entropy of the Rainy dataset: 0.972

```
In [11]: entropy_and_infogain(rainy, 'temp')
```

```
temp play
3 mild yes
9 mild yes
13 mild no
Entropy of temp - mild = 0.918

temp play
4 cool yes
5 cool no
Entropy of temp - cool = 1.8
Information Gain for temp = 0.62
```

```
In [12]: entropy_and_infogain(rainy, 'humidity')
```

```
humidity play
3 high yes
13 high no
Entropy of humidity - high = 1.0

humidity play
4 normal yes
5 normal no
9 normal yes
Entropy of humidity - normal = 0.918
Information Gain for humidity = 0.62
```

```
In [13]: entropy_and_infogain(rainy, 'windy')
```

```
windy play
3 False yes
4 False yes
9 False yes
Entropy of windy - False = 0.0

windy play
5 True no
13 True no
Entropy of windy - True = 0.0
Information Gain for windy = 0.972
```

wind has highest information gain

Output

Output:

Entropy of the dataset : 0.9331

Pregnancies - Entropy : 3.482, IG : 0.062

Glucose - Entropy : 6.751, IG : 0.304

BloodPressure - Entropy : 4.792, IG : 0.059

SkinThickness - Entropy : 4.586, IG : 0.022

Insulin - Entropy : 4.682, IG : 0.277

BMI - Entropy : 7.594, IG : 0.344

DiabetesPedigreeFunction - Entropy : 8.829, IG : 0.65

Age - Entropy : 5.029, IG : 0.141

PROGRAM 4

Date:19-04-2024

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

LINEAR REGRESSION:

Algorithm

cs/05/2024

Program-4

Implement Linear and Multi-Linear Regression algorithm.

Linear Regression:

```
function linear_regression(X, y, learning_rate, num_iterations)
    Initialize random values for slope (m)
    & intercept (b)

    for i = 1 to num_iterations:

        predictions = m * X + b

        errors = predictions - y

        loss = mean_squared_error(errors)

        gradient_m = (2/N) * sum(errors * X)
        gradient_b = (2/N) * sum(errors)
        m = m - learning_rate * gradient_m
        b = b - learning_rate * gradient_b

    Return m, b

function mean_squared_error(errors):
    squared_errors = errors^2
    mse = sum(squared_errors) / sum(errors)
    return mse
```

Code

Importing Dataset

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [18]: df=pd.read_csv("/Salary_Data.csv")
df
```

```
Out[18]:
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93640.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112693.0
28	10.3	122391.0
29	10.5	121872.0

Slope and Intercept calculation

```
In [19]: def linear(X, b0, b1):  
        return [b0+b1*x for x in X]  
  
In [20]: # b0 - Intercept  
        def intercept(X, Y, b1):  
            x_ = np.mean(X)  
            y_ = np.mean(Y)  
  
            return y_-b1*x_  
  
In [21]: # b1 - Slope  
        def slope(X, Y):  
            x_ = np.mean(X)  
            y_ = np.mean(Y)  
  
            rise = sum([(x-x_) * (y-y_) for x,y in zip(X,Y)])  
            run = sum([(x-x_)**2 for x,y in zip(X,Y)])  
  
            return rise / run
```

Predicted Values Graph

```
In [17]: plt.figure(figsize=(8,5))  
        plt.title("YearsExperience vs Salary")  
        plt.scatter(predictor, target, color = "#4682B4")  
        plt.xlabel("YearsExperience")  
        plt.ylabel("Salary")  
        plt.show()
```




```
In [18]: b1 = slope(predictor, target)  
        b0 = intercept(predictor, target, b1)  
        predicted = linear(predictor, b0, b1)
```

```
In [18]: plt.figure(figsize = (8, 5))  
        plt.plot(predictor, predicted, color = "#D2691E")  
        plt.scatter(predictor, predicted, color = "#D2691E"). #red line is predicted data  
        plt.title("Predicted values by Linear Regression", fontsize = 15)  
        plt.xlabel("YearsExperience")  
        plt.ylabel("Salary")  
        plt.scatter(predictor, target, color = "#4682B4")  
        plt.show()
```

Output

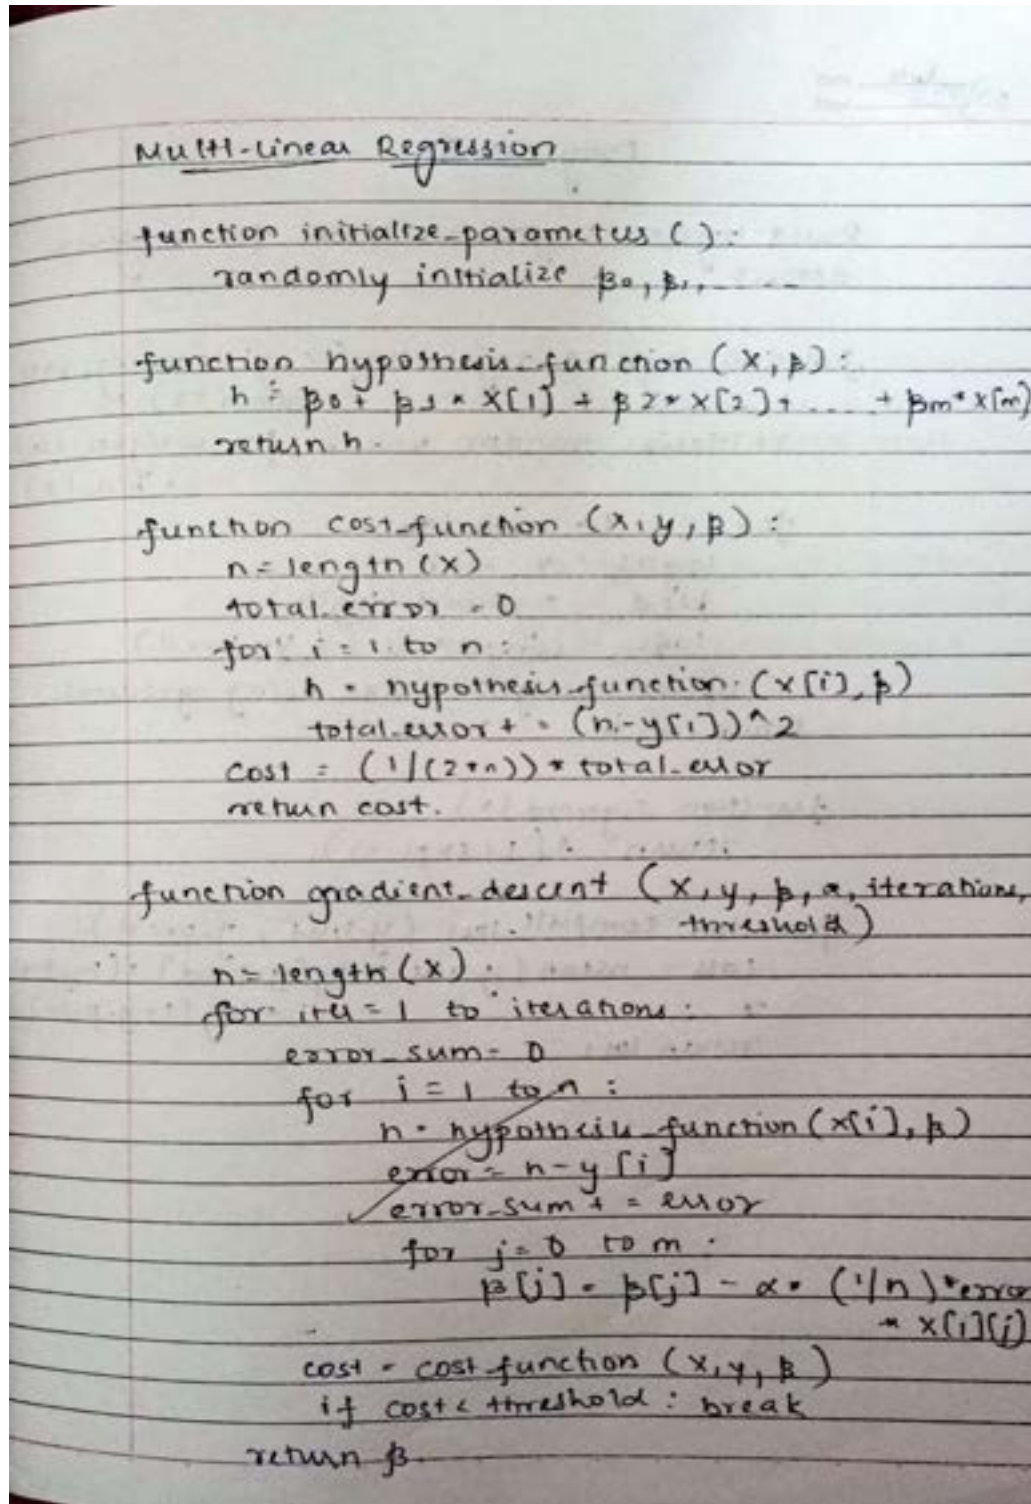
```
In [28]: print("Coefficients:\n\n")
print("b0 : ", b0)
print("b1 : ", b1)
```

```
Coefficients:
\n\n
b0 : 25792.20019866669
b1 : 9449.962321455077
```



MULTIPLE LINEAR REGRESSION:

Algorithm



```
Multi-linear Regression

function initialize-parameters():
    randomly initialize  $\beta_0, \beta_1, \dots$ 

function hypothesis-function( $X, \beta$ ):
     $h = \beta_0 + \beta_1 * X[1] + \beta_2 * X[2] + \dots + \beta_m * X[m]$ 
    return h

function cost-function( $X, y, \beta$ ):
     $n = \text{length}(X)$ 
    total-error = 0
    for  $i = 1$  to  $n$ :
         $h = \text{hypothesis-function}(X[i], \beta)$ 
        total-error +=  $(h - y[i])^2$ 
    cost =  $(1/(2*n)) * \text{total-error}$ 
    return cost

function gradient-descent( $X, y, \beta, \alpha, \text{iterations}, \text{threshold}$ ):
     $n = \text{length}(X)$ 
    for  $it = 1$  to  $\text{iterations}$ :
        error-sum = 0
        for  $i = 1$  to  $n$ :
             $h = \text{hypothesis-function}(X[i], \beta)$ 
            error =  $h - y[i]$ 
            error-sum += error
        for  $j = 0$  to  $m$ :
             $\beta[j] = \beta[j] - \alpha * (1/n) * \text{error} * X[i][j]$ 
        cost = cost-function( $X, y, \beta$ )
        if cost < threshold: break
    return  $\beta$ 
```


Code

```
In [18]: house = pd.read_csv('https://github.com/YBIFoundation/Dataset/raw/main/Boston.csv')

In [19]: house.head()

Out[19]:
```

	CRIM	ZN	INDUS	CHAS	NX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	5.33	36.2

```
In [21]: house.describe()

Out[21]:
```

	CRIM	ZN	INDUS	CHAS	NX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554895	6.284534	68.574901	3.796043	9.549407	408.237154	18.455511	396.900000	15.230214	32.047171
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.537116	2.164680	158.504016	4.543417	4.743417
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000	12.600000	392.830000	4.030000	21.600000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.000000	17.400000	392.830000	4.030000	21.600000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.206500	77.500000	3.207450	5.000000	330.000000	19.050000	396.900000	9.140000	21.600000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.000000	20.200000	396.900000	9.140000	21.600000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000	9.140000	21.600000

◀ ▶

```
In [22]: house.columns

Out[22]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
          'PTRATIO', 'B', 'LSTAT', 'MEDV'],  
         dtype='object')
```

```
In [23]: y = house["MEDV"]

In [24]: X = house.drop(["MEDV"],axis=1)

In [25]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X,y, train_size=0.7, random_state=2529)

In [26]: X_train.shape, X_test.shape, y_train.shape, y_test.shape

Out[26]: (354, 13), (152, 13), (354, 1), (152, 1)
```

```

In [25]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, random_state=2529)

In [26]: X_train.shape, X_test.shape, y_train.shape, y_test.shape

Out[26]: ((354, 13), (152, 13), (354,), (152,))

In [27]: from sklearn.linear_model import LinearRegression
model = LinearRegression()

In [28]: # Step 6 : train or fit model
model.fit(X_train, y_train)

Out[28]: LinearRegression()
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [29]: model.intercept_

Out[29]: 34.21916308862993

In [30]: model.coef_

Out[30]: array([-1.29e+01,  3.05e+02,  1.54e+02,  2.35e+00, -2.04e+01,  4.41e+00,
        4.61e+03, -1.59e+00,  2.51e+01, -9.60e+03, -9.64e+01,  1.01e+02,
        -5.43e+01])

In [31]: # Step 7 : predict model
y_pred = model.predict(X_test)

```

Output

```

In [32]: y_pred

Out[32]: array([31.72, 22.02, 21.17, 39.70, 20.1 , 22.86, 18.36, 14.79, 22.56,
        21.35, 18.38, 27.97, 29.86,  6.45, 10.68, 26.25, 23.89, 25.23,
        3.62, 36.22, 24.00, 22.94, 14.27, 20.79, 24.23, 16.74, 18.75,
        20.97, 28.51, 20.86,  9.23, 17.07, 22.07, 22.23, 39.26, 26.17,
        42.5 , 19.35, 34.52, 14.07, 13.81, 23.28, 13.79,  9.01, 23.65,
        25.55, 18.17, 16.82, 14.86, 14.86, 33.70, 33.27, 15.49, 24.00,
        27.64, 19.50, 45.02, 20.97, 20.07, 27.67, 34.59, 12.71, 23.66,
        31.66, 28.97, 32.46, 13.93, 35.49, 19.36, 19.6 ,  1.44, 24.1 ,
        33.67, 20.62, 26.89, 21.29, 31.95, 29.74, 13.93, 13.82, 19.76,
        21.54, 20.87, 23.63, 28.8 , 23.64,  6.95, 22.2 , -6.82, 16.97,
        16.77, 25.44, 14.95,  3.72, 15.03, 16.91, 21.46, 31.66, 30.72,
        23.73, 22.19, 13.76, 18.47, 18.15, 36.6 , 27.49, 11. , 17.26,
        22.49, 16.53, 29.49, 22.89, 24.60, 20.30, 19.60, 22.55, 27.32,
        24.86, 20.2 , 29.14,  7.43,  5.85, 25.35, 38.73, 23.94, 25.20,
        20.11, 19.75, 25.07, 35.16, 27.32, 27.26, 31.4 , 16.55, 14.3 ,
        23.77,  7.65, 23.35, 21.37, 26.12, 25.32, 13.12, 17.67, 36.2 ,
        20.5 , 27.95, 22.46, 18.15, 31.24, 20.85, 27.36, 30.53])

In [33]: # Step 8 : model accuracy
from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error, mean_squared_error

In [34]: mean_absolute_error(y_test, y_pred)

Out[34]: 3.155030927602485

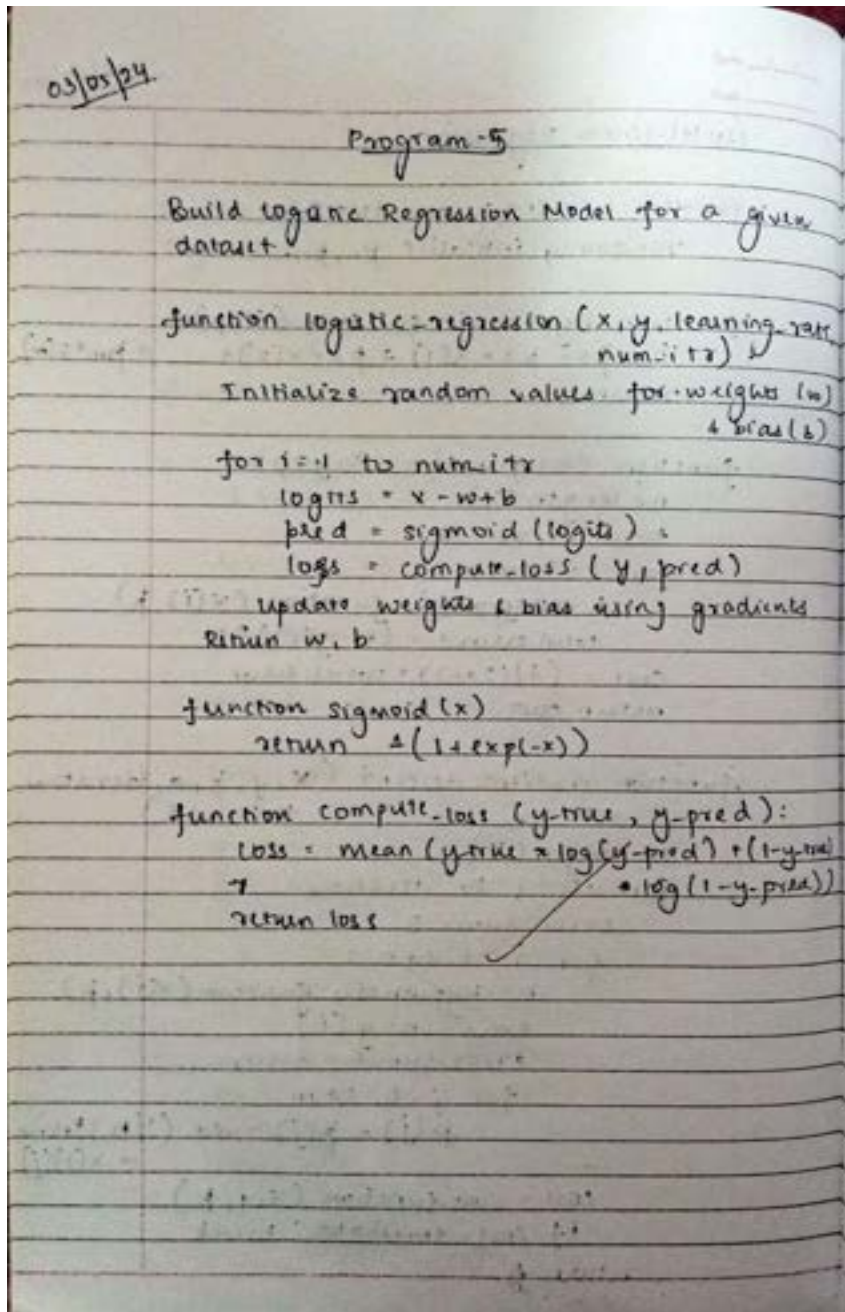
```

PROGRAM 5

Date:03-05-2024

Build Logistic Regression Model for a given dataset

Algorithm



Code

```
In [2]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list the files in the input directory

import os


In [4]: data = pd.read_csv("../content/data.csv")


In [5]: data.drop(['Unnamed: 32', "id"], axis=1, inplace=True)
data.diagnosis = [1 if each == "W" else 0 for each in data.diagnosis]
y = data.diagnosis.values
x_data = data.drop(['diagnosis'], axis=1)


In [7]: # Assuming x_data is a numpy array or pandas Dataframe
x = (x_data - np.min(x_data)) / (np.max(x_data) - np.min(x_data))


In [8]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.15, random_state=42)

x_train = x_train.T
x_test = x_test.T
y_train = y_train.T
y_test = y_test.T

print("x train: ", x_train.shape)
print("x test: ", x_test.shape)
print("y train: ", y_train.shape)
print("y test: ", y_test.shape)

x_train: (30, 483)
x_test: (30, 86)
y_train: (483,)
y_test: (86,)


In [9]: def initialize_weights_and_bias(dimension):
    w = np.full((dimension,1),0.01)
    b = 0.0
    return w, b


In [10]: def sigmoid(z):
    y_head = 1/(1+np.exp(-z))
    return y_head


In [ ]: def forward_backward_propagation(w,b,x_train,y_train):
    # forward propagation
    z = np.dot(w.T,x_train) + b
    y_head = sigmoid(z)
    loss = -y_train*np.log(y_head)-(1-y_train)*np.log(1-y_head)
    cost = (np.sum(loss))/x_train.shape[1] # x_train.shape[1] is for scaling
    # backward propagation
    derivative_weight = (np.dot(x_train,((y_head-y_train).T)))/x_train.shape[1] # x_train.shape[1] is for scaling
    derivative_bias = np.sum(y_head-y_train)/x_train.shape[1] # x_train.shape[1] is for scaling
    gradients = {"derivative_weight": derivative_weight, "derivative_bias": derivative_bias}
    return cost,gradients
```

```
In [1]: def update(u, b, x_train, y_train, learning_rate, number_of_iteration):
cost_list = []
cost_list2 = []
index = []
# updating(learning) parameters (x number_of_iteration times)
for i in range(number_of_iteration):
    # make forward and backward propagation and find cost and gradients
    cost, gradients = forward_backward_propagation(u, b, x_train, y_train)
    cost_list.append(cost)
    u = u - learning_rate * gradients["derivative_weight"]
    b = b - learning_rate * gradients["derivative_bias"]
    if i % 10 == 0:
        cost_list2.append(cost)
        index.append(i)
        print ("Cost after iteration %i: %f" % (i, cost))
# we update(learn) parameters weights and bias
parameters = {"weight": u, "bias": b}
plt.plot(index, cost_list2)
plt.xticks(index, rotation='vertical')
plt.xlabel("Number of Iteration")
plt.ylabel("Cost")
plt.show()
return parameters, gradients, cost_list
```

```
In [11]: def predict(u, b, x_test):
# x_test is a input for forward propagation
z = sigmoid(np.dot(u.T, x_test) + b)
Y_prediction = np.zeros([1, x_test.shape[1]])
# if z is bigger than 0.5, our prediction is sign one (y_head=1),
# if z is smaller than 0.5, our prediction is sign zero (y_head=0),
for i in range(x_test.shape[1]):
    if z[0,i] <= 0.5:
        Y_prediction[0,i] = 0
    else:
        Y_prediction[0,i] = 1
return Y_prediction
```

```
In [18]: def sigmoid(z):
return 1 / (1 + np.exp(-z))

def initialize_weights_and_bias(dim):
u = np.zeros((dim, 1))
b = 0
return u, b

def compute_cost(u, b, x, y):
m = x.shape[1]
A = sigmoid(np.dot(u.T, x) + b)
cost = -1 / m * np.sum(y * np.log(A) + (1 - y) * np.log(1 - A))
return cost

def propagate(u, b, x, y):
m = x.shape[1]
A = sigmoid(np.dot(u.T, x) + b)
du = 1 / m * np.dot(x, (A - y).T)
db = 1 / m * np.sum(A - y)
return du, db
```

```

def logistic_regression(x_train, y_train, x_test, y_test, learning_rate, num_iterations):
    # Initialize
    dimension = x_train.shape[0] # Number of features
    w, b = initialize_weights_and_bias(dimension)
    costs = []

    # Gradient Descent
    for i in range(num_iterations):
        # Forward and Backward Propagation
        dw, db = propagate(w, b, x_train, y_train)

        # Update parameters
        w -= learning_rate * dw
        b -= learning_rate * db

        # Record the costs
        if i % 100 == 0:
            cost = compute_cost(w, b, x_train, y_train)
            costs.append(cost)
            print(f"Cost after iteration {i}: {cost}")

    # Evaluate model
    y_prediction_train = predict(w, b, x_train)
    y_prediction_test = predict(w, b, x_test)

    train_accuracy = 100 - np.mean(np.abs(y_prediction_train - y_train)) * 100
    test_accuracy = 100 - np.mean(np.abs(y_prediction_test - y_test)) * 100

    print("Train accuracy: {} %".format(train_accuracy))
    print("Test accuracy: {} %".format(test_accuracy))

    return w, b

# Assuming you have defined the predict function
# def predict(w, b, x):
#     ...

# Assuming you have defined x_train, y_train, x_test, y_test, learning_rate, and num_iterations
logistic_regression(x_train, y_train, x_test, y_test, learning_rate=1, num_iterations=100)

```

Output

```
Cost after iteration 0: 0.6782740160052536
Train accuracy: 80.74534161490683 %
Test accuracy: 81.3953488372093 %
```

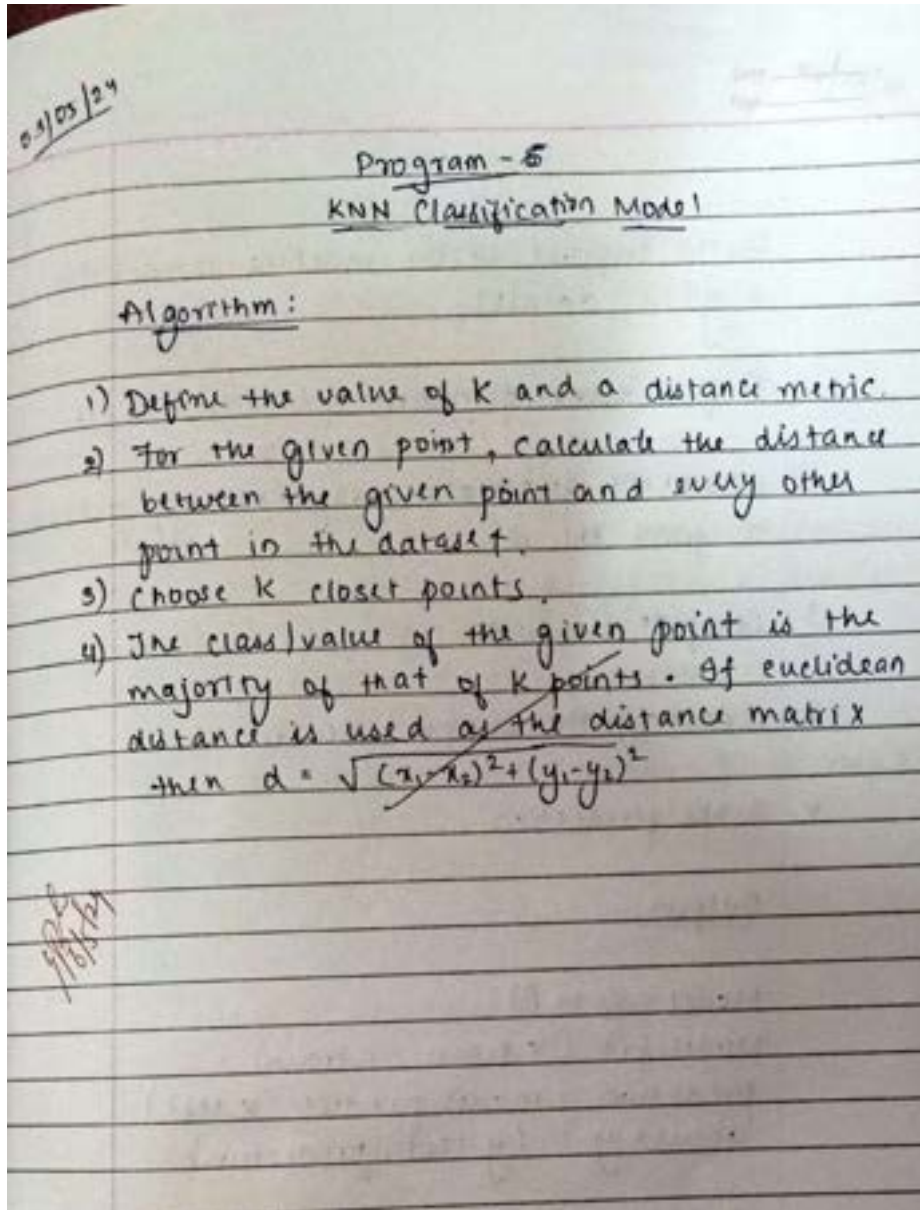
```
Out[18]: (array([[ 1.77806654e-02],
 [ 1.10160388e-02],
 [ 1.27806976e-01],
 [ 1.95749649e+00],
 [ 1.85931875e-05],
 [ 2.68863405e-04],
 [ 4.89020048e-04],
 [ 2.63106803e-04],
 [ 3.49357933e-05],
 [-2.02145931e-05],
 [ 1.25690784e-03],
 [-3.98285024e-04],
 [ 8.96937014e-03],
 [ 2.02426962e-01],
 [-3.60718647e-06],
 [ 4.19150446e-05],
 [ 6.03411729e-05],
 [ 2.00740406e-05],
 [-6.24803672e-06],
 [ 6.24944780e-07],
 [ 2.79506973e-02],
 [ 1.99326360e-02],
 [ 1.98774929e-01],
 [ 3.39189908e+00],
 [ 5.79135019e-05],
 [ 8.53041205e-04],
 [ 1.25862280e-03],
 [ 4.60695564e-04],
 [ 1.89671301e-04],
 [ 3.52490835e-05]]),
-1.5161875221606185)
```


PROGRAM 6

Date:19-04-2024

Build KNN Classification model for a given dataset.

Algorithm



Code

```
In [1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # for data visualization purposes
import seaborn as sns # for data visualization
%matplotlib inline

In [2]: data = '/content/cancer_detector.txt'
df = pd.read_csv(data, header=None)

In [3]: df.shape

Out[3]: (699, 11)

In [4]: col_names = ['Id', 'Clump_thickness', 'Uniformity_Cell_Size', 'Uniformity_Cell_Shape', 'Marginal_Adhesion',
                    'Single_Epithelial_Cell_Size', 'Bare_Nuclei', 'Bland_Chromatin', 'Normal_Nucleoli', 'Mitoses', 'Class']

df.columns = col_names

df.columns

Out[4]: Index(['Id', 'Clump_thickness', 'Uniformity_Cell_Size',
              'Uniformity_Cell_Shape', 'Marginal_Adhesion',
              'Single_Epithelial_Cell_Size', 'Bare_Nuclei', 'Bland_Chromatin',
              'Normal_Nucleoli', 'Mitoses', 'Class'],
              dtype='object')

In [5]: df.head()

Out[5]:
```

	Id	Clump_thickness	Uniformity_Cell_Size	Uniformity_Cell_Shape	Marginal_Adhesion	Single_Epithelial_Cell_Size	Bare_Nuclei	Bla
0	1000025	5	1	1	1	2	1	
1	1002945	5	4	4	5	7	10	
2	1015425	3	1	1	1	2	2	
3	1016277	6	8	8	1	3	4	
4	1017023	4	1	1	3	2	1	

4

```
In [10]: import numpy as np
```

```
In [11]: # view summary statistics for numerical variables
```

```
print(round(df.describe(),2))
```

```

Clump_thickness  Uniformity_Cell_Size  Uniformity_Cell_Shape  \
count      699.00                699.00                699.00
mean         9.42                  3.13                  3.21
std          2.82                  1.08                  1.07
min          1.00                  1.00                  1.00
25%          2.00                  1.00                  1.00
50%          4.00                  1.00                  1.00
75%          6.00                  3.00                  3.00
max         19.00                 10.00                 10.00

```

```

Marginal_Adhesion  Single_Epithelial_Cell_Size  Bare_Nuclei  \
count      699.00                699.00                699.00
mean         2.81                  3.22                  3.54
std          2.06                  1.21                  3.64
min          1.00                  1.00                  1.00
25%          1.00                  2.00                  1.00
50%          1.00                  3.00                  1.00
75%          4.00                  4.00                  4.00
max         10.00                 10.00                 10.00

```

```

Gland_Chromatin  Normal_Nucleoli  Mitoses  Class
count      699.00                699.00                699.00
mean         1.44                  2.87                  3.09
std          1.44                  3.05                  1.72
min          1.00                  1.00                  1.00
25%          1.00                  1.00                  1.00
50%          1.00                  1.00                  1.00
75%          3.00                  4.00                  4.00
max         10.00                 10.00                 10.00

```

```
In [12]: x = df.drop(['Class'], axis=1)
```

```
y = df['Class']
```

```
In [13]: from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state = 0)
```

```
In [14]: x_train.shape, x_test.shape
```

```
Out[14]: ((559, 9), (140, 9))
```

```
In [25]: for col in X_train.columns:
         if X_train[col].isnull().mean()>0:
             print(col, round(X_train[col].isnull().mean(),4))
```

```
Bare_Nuclei 0.0233
```

```
In [16]: for df1 in [X_train, X_test]:
         for col in X_train.columns:
             col_median=X_train[col].median()
             df1[col].fillna(col_median, inplace=True)
```

```
In [27]: cols = X_train.columns
```

```
In [18]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
In [20]: X_train = pd.DataFrame(X_train, columns=cols)
```

```
In [30]: X_test = pd.DataFrame(X_test, columns=cols)
```

Output

```
In [33]: y_pred = knn.predict(X_test)
         y_pred
```

```
Out[33]: array([2, 2, 4, 2, 4, 2, 4, 2, 4, 2, 2, 2, 4, 4, 4, 2, 2, 4, 4, 2, 4, 4,
                2, 2, 4, 2, 2, 4, 4, 2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2,
                4, 4, 2, 4, 2, 4, 4, 2, 2, 4, 2, 2, 2, 2, 2, 2, 4, 2, 2, 4, 4, 4,
                4, 2, 2, 4, 2, 2, 4, 4, 2, 2, 2, 4, 2, 2, 2, 4, 2, 2, 2, 4, 2,
                4, 4, 2, 2, 2, 4, 2, 2, 2, 4, 2, 4, 4, 2, 2, 2, 4, 2, 2, 2, 2,
                4, 4, 4, 2, 2, 2, 2, 4, 4, 4, 4, 2, 4, 2, 2, 4, 4, 4, 4, 4, 2,
                2, 4, 4, 2, 2, 4, 2, 2])
```

```
In [34]: knn.predict_proba(X_test)[:,0]
```

```
Out[34]: array([[1.         , 1.         , 0.33333333, 1.         , 0.         ,
                 1.         , 0.         , 1.         , 0.         , 0.66666667,
                 1.         , 1.         , 0.         , 0.33333333, 0.         ,
                 1.         , 1.         , 0.         , 0.         , 1.         ,
                 0.         , 0.         , 1.         , 1.         , 1.         ,
                 0.         , 1.         , 1.         , 1.         , 1.         ,
                 0.66666667, 1.         , 0.         , 1.         , 1.         ,
                 1.         , 1.         , 1.         , 1.         , 1.         ,
                 0.         , 1.         , 1.         , 1.         , 1.         ,
                 1.         , 1.         , 1.         , 0.66666667, 1.         ,
                 0.         , 1.         , 1.         , 0.         , 0.         ,
                 0.33333333, 0.         , 1.         , 1.         , 0.         ,
                 1.         , 1.         , 0.         , 0.         , 1.         ,
                 1.         , 1.         , 1.         , 1.         , 1.         ,
                 1.         , 1.         , 1.         , 1.         , 1.         ,
                 0.         , 1.         , 1.         , 1.         , 1.         ,
                 0.         , 0.33333333, 1.         , 0.         , 1.         ,
                 1.         , 0.33333333, 0.33333333, 0.         , 0.         ,
                 0.         , 1.         , 1.         , 0.33333333, 0.         ,
                 1.         , 1.         , 0.         , 1.         , 1.         ])
```

```
In [35]: from sklearn.metrics import accuracy_score
         print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_pred)))
```

Model accuracy score: 0.9714

```
In [36]: y_pred_train = knn.predict(X_train)
```

```
In [37]: print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pred_train)))
```

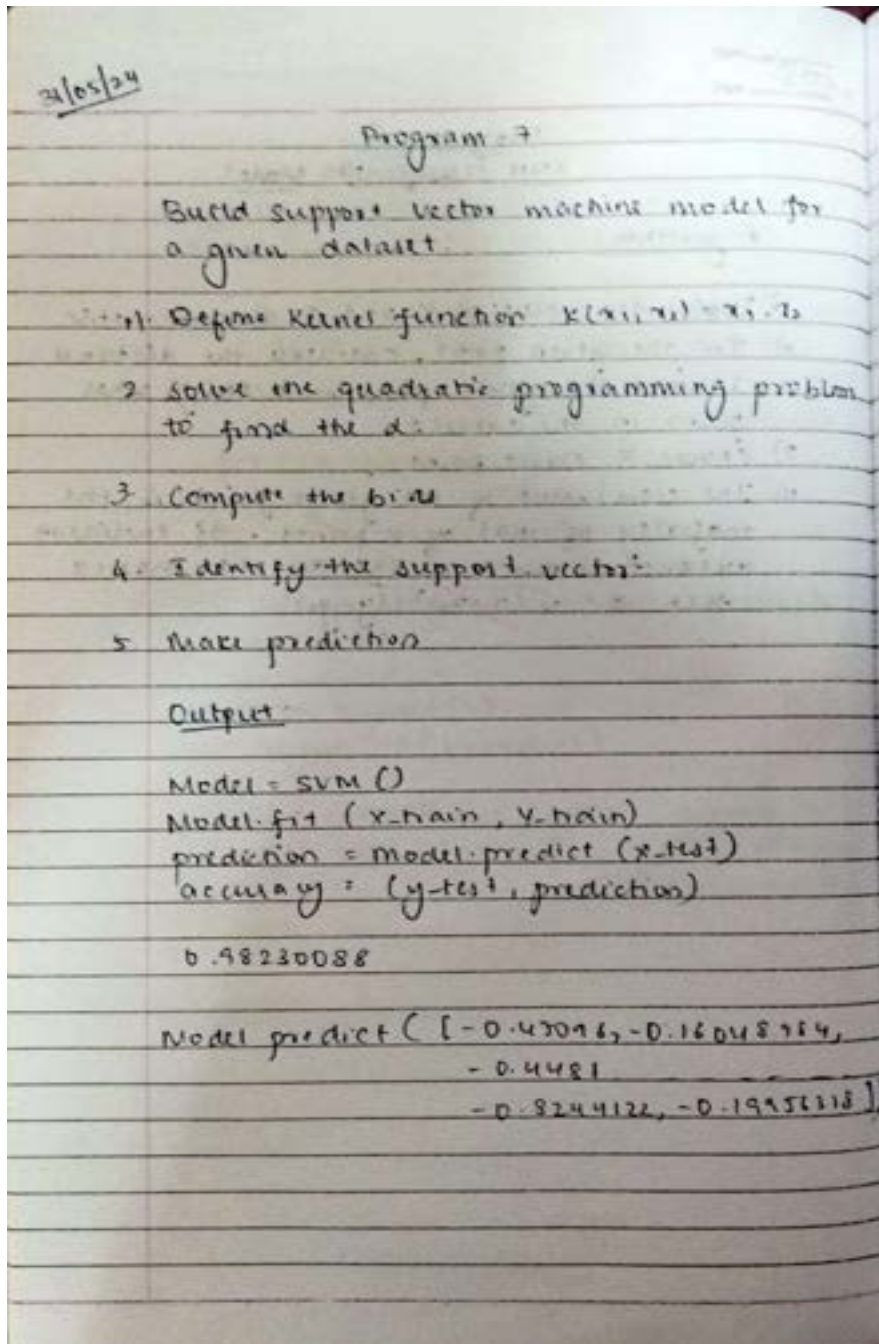
Training-set accuracy score: 0.9811

PROGRAM 7

Date:24-05-2024

Build Support vector machine model for a given dataset

Algorithm



Code

```

import seaborn as sns
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px

df = pd.read_csv("../content/breast-cancer.csv")
df.head()

df

```

	id	diagnosis	radius.mean	texture.mean	perimeter.mean	area.mean	smoothness.mean	compactness.mean	concavity.mean
0	84232	M	17.99	10.38	122.80	1327.0	0.11840	0.27760	0.840
1	84257	M	20.57	17.77	152.90	1526.0	0.08475	0.07864	0.086
2	843080	M	19.69	21.25	780.00	1229.0	0.12960	0.10980	0.165
3	8434001	M	11.42	20.34	77.58	586.1	0.14250	0.28100	0.245
4	843842	M	20.29	14.34	185.10	1297.0	0.10000	0.11280	0.198

5 rows x 10 columns

```

df.drop("id", axis=1, inplace=True) #drop redundant column
df.describe().T

```

	count	mean	std	min	25%	50%	75%	max
radius.mean	569.0	14.12752	5.524349	6.981000	11.700000	15.370000	18.760000	28.110000
texture.mean	569.0	16.28649	4.301096	5.710000	14.770000	18.340000	21.800000	35.280000
perimeter.mean	569.0	91.96303	24.298901	43.700000	75.700000	86.240000	104.100000	188.800000
area.mean	569.0	654.88104	331.814129	143.800000	420.400000	555.100000	782.700000	2701.000000
smoothness.mean	569.0	0.096480	0.014364	0.052630	0.060370	0.069870	0.101400	0.16340
compactness.mean	569.0	0.104341	0.052818	0.019980	0.044920	0.076260	0.130400	0.34540
concavity.mean	569.0	0.087799	0.075720	0.000000	0.029660	0.067540	0.130700	0.42680
concave points.mean	569.0	0.048919	0.038803	0.000000	0.020100	0.033000	0.074000	0.20120
symmetry.mean	569.0	0.181142	0.027414	0.106000	0.141900	0.176200	0.195700	0.35400
fractal dimension.mean	569.0	0.062798	0.007060	0.049960	0.057780	0.067540	0.066120	0.08744
radius.se	569.0	0.405772	0.277118	0.111800	0.252400	0.324200	0.478900	1.07000
texture.se	569.0	1.216952	0.951948	0.360200	0.831900	1.108000	1.474000	4.88000
perimeter.se	569.0	2.860359	2.521955	0.757000	1.606000	2.287000	3.357000	21.98000
area.se	569.0	40.507079	45.491006	4.852000	17.800000	24.300000	45.190000	542.20000
smoothness.se	569.0	0.007041	0.000309	0.001719	0.001409	0.000380	0.001146	0.00118
compactness.se	569.0	0.025478	0.017808	0.002252	0.010880	0.020400	0.030480	0.10540
concavity.se	569.0	0.013894	0.001186	0.000000	0.010090	0.023890	0.042000	0.19600
concave points.se	569.0	0.011796	0.006170	0.000000	0.007628	0.010960	0.014710	0.06278
symmetry.se	569.0	0.020542	0.008206	0.007962	0.018180	0.018730	0.023480	0.07896
fractal dimension.se	569.0	0.003785	0.002646	0.000000	0.002248	0.001187	0.004038	0.02984
radius.min	569.0	16.269190	4.816242	7.900000	14.010000	14.970000	16.790000	86.34000
texture.min	569.0	25.677223	8.146258	12.000000	21.000000	25.410000	29.720000	49.54000
perimeter.min	569.0	107.261218	19.402542	35.410000	64.100000	87.660000	125.400000	251.20000
area.min	569.0	880.581128	940.036999	185.200000	515.300000	686.500000	1084.000000	4254.00000
smoothness.min	569.0	0.152369	0.022852	0.071170	0.119600	0.111000	0.146000	0.22260
compactness.min	569.0	0.254208	0.107106	0.070290	0.147200	0.211900	0.339100	1.06000
concavity.min	569.0	0.272188	0.208624	0.000000	0.154000	0.226700	0.367900	1.25200
concave points.min	569.0	0.114606	0.008752	0.000000	0.064900	0.099900	0.161400	0.29100
symmetry.min	569.0	0.293075	0.041867	0.156000	0.250400	0.282200	0.317900	0.66080
fractal dimension.min	569.0	0.001946	0.018803	0.000040	0.007460	0.000040	0.002080	0.00750

```

30 | | df['diagnosis'] = df['diagnosis'].astype(int) #convert the label into 1/0
31 | |
32 | |
33 | | # Get the absolute value of the correlation
34 | | cor_target = abs(cor['diagnosis'])
35 | |
36 | | # Select highly correlated features (threshold = 0.2)
37 | | relevant_features = cor_target[cor_target>0.2]
38 | |
39 | | # Collect the names of the features
40 | | names = [index for index, value in relevant_features.items()]
41 | |
42 | | # Drop the target variable from the results
43 | | names.remove('diagnosis')
44 | |
45 | | # Display the results
46 | | print(names)
47 | |
48 | | ['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean', 'concave points_mean', 'symmetry_mean', 'fractal_dimension', 'radius_w', 'perimeter_w', 'area_w', 'compactness_w', 'concavity_w', 'concave points_w', 'symmetry_w', 'fractal_dimension_w']
49 | |
50 | | X = df[names].values
51 | | y = df['diagnosis']
52 | |
53 | |
54 | | def scale(X):
55 | |     """
56 | |     Standardizes the data in the array X.
57 | |
58 | |     Parameters:
59 | |     X (numpy.ndarray): Features array of shape (n_samples, n_features).
60 | |
61 | |     Returns:
62 | |     numpy.ndarray: The standardized features array.
63 | |     """
64 | |     # Calculate the mean and standard deviation of each feature
65 | |     mean = np.mean(X, axis=0)
66 | |     std = np.std(X, axis=0)
67 | |
68 | |     # Standardize the data
69 | |     X = (X - mean) / std
70 | |     return X
71 | |
72 | | X = scale(X)
73 | |
74 | |
75 | | def train_test_split(X, y, random_state=0, test_size=0.2):
76 | |     """
77 | |     Splits the data into training and testing sets.
78 | |
79 | |     Parameters:
80 | |     X (numpy.ndarray): Features array of shape (n_samples, n_features).
81 | |     y (numpy.ndarray): Target array of shape (n_samples,).
82 | |     random_state (int): Seed for the random number generator. Default is 0.
83 | |     test_size (float): Proportion of samples to include in the test set. Default is 0.2.
84 | |
85 | |     Returns:
86 | |     tuple[numpy.ndarray]: A tuple containing X_train, X_test, y_train, y_test.
87 | |     """
88 | |     # Get number of samples
89 | |     n_samples = X.shape[0]
90 | |
91 | |     # Set the seed for the random number generator
92 | |     np.random.seed(random_state)
93 | |
94 | |     # Shuffle the indices
95 | |     shuffled_indices = np.random.permutation(np.arange(n_samples))
96 | |
97 | |     # Determine the size of the test set
98 | |     test_size = int(n_samples * test_size)
99 | |
100 | |     # Split the indices into test and train
101 | |     test_indices = shuffled_indices[test_size:]
102 | |     train_indices = shuffled_indices[:test_size]
103 | |
104 | |     # Split the features and target arrays into test and train
105 | |     X_train, X_test = X[train_indices], X[test_indices]
106 | |     y_train, y_test = y[train_indices], y[test_indices]
107 | |
108 | |     return X_train, X_test, y_train, y_test
109 | |
110 | | X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=0) #split the data into training and testing sets

```

```

In [ ]: class SVM:

    def __init__(self, iterations=1000, lr=0.01, lambdas=0.01):

        self.lambdas = lambdas
        self.iterations = iterations
        self.lr = lr
        self.w = None
        self.b = None
    def initialize_parameters(self, X):

        m, n = X.shape
        self.w = np.zeros(n)
        self.b = 0

    def gradient_descent(self, X, y):

        y_ = np.where(y <= 0, -1, 1)
        for i, x in enumerate(X):
            if y_[i] * (np.dot(x, self.w) - self.b) >= 1:
                dw = 2 * self.lambdas * self.w
                db = 0
            else:
                dw = 2 * self.lambdas * self.w - np.dot(x, y_[i])
                db = y_[i]
            self.update_parameters(dw, db)

    def update_parameters(self, dw, db):

        self.w = self.w - self.lr * dw
        self.b = self.b - self.lr * db
    def fit(self, X, y):

        self.initialize_parameters(X)
        for i in range(self.iterations):
            self.gradient_descent(X, y)

    def predict(self, X):

        # get the outputs
        output = np.dot(X, self.w) - self.b
        # get the signs of the labels depending on if it's greater/less than zero
        label_signs = np.sign(output)
        # set predictions to 0 if they are less than or equal to -1 else set them to 1
        predictions = np.where(label_signs <= -1, 0, 1)
        return predictions

In [ ]: def accuracy(y_true, y_pred):

    total_samples = len(y_true)
    correct_predictions = np.sum(y_true == y_pred)
    return (correct_predictions / total_samples)

```

Output

```

In [ ]: model = SVM()
        model.fit(X_train, y_train)
        predictions = model.predict(X_test)
        accuracy(y_test, predictions)

```

```

Out[ ]: 0.9823088849557522

```


PROGRAM 8

Date: 31-05-2024

Build Artificial Neural Network model with back propagation on a given dataset

Algorithm

31/5/24

Program 8

Build Artificial Neural Network model with back propagation on a given dataset.

Algorithm:

1. Create a feed-forward network with n_i inputs, n_{hidden} hidden units, and n_{out} output units.
2. Initialize all network weights to small random numbers.
3. Until the termination condition is met, do
 - For each (x, E) , in training examples, do
 - Propagate the input forward through the network:
 - 1. Input the instance x to the network and compute the output o_k of every unit k in the network.
 - Propagate the error backward through the network:
 - 2. For each network output unit k , calculate its error term δ_k
$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$
 - 3. For each hidden unit h , calculate its error term δ_h
$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{hk} \delta_k$$
 - 4. Update each network weight w_{ji}
$$w_{ji} \leftarrow w_{ji} + \Delta w_{ji}$$
where
$$\Delta w_{ji} = \eta \delta_j x_i$$

Output: Learning Accuracy = 0.66233, Testing Accuracy = 0

Code

```
In [1]: import numpy as np
        from sklearn.model_selection import train_test_split

        db = np.loadtxt("../dataset/dmn-breast-cancer.txt")
        print("Database raw shape (%s,%s)" % np.shape(db))

Database raw shape (86,7159)

In [2]: np.random.shuffle(db)
        y = db[:, 0]
        x = np.delete(db, [0], axis=1)
        x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.1)
        print(np.shape(x_train), np.shape(x_test))

(77, 7129) (9, 7129)

In [3]: hidden_layer = np.zeros(72)
        weights = np.random.random((len(x[0]), 72))
        output_layer = np.zeros(2)
        hidden_weights = np.random.random((72, 2))

In [4]: def sum_function(weights, index_locked_col, x):
        result = 0
        for i in range(0, len(x)):
            result += x[i] * weights[i][index_locked_col]
        return result

In [5]: def activate_layer(layer, weights, x):
        for i in range(0, len(layer)):
            layer[i] = 1.7159 * np.tanh(2.0 * sum_function(weights, i, x) / 3.0)

In [6]: def soft_max(layer):
        soft_max_output_layer = np.zeros(len(layer))
        for i in range(0, len(layer)):
            denominator = 0
            for j in range(0, len(layer)):
                denominator += np.exp(layer[j] + np.max(layer))
            soft_max_output_layer[i] = np.exp(layer[i] + np.max(layer)) / denominator
        return soft_max_output_layer

In [7]: def recalculate_weights(learning_rate, weights, gradient, activation):
        for i in range(0, len(weights)):
            for j in range(0, len(weights[i])):
                weights[i][j] = (learning_rate * gradient[j] * activation[i] + weights[i][j])

In [8]: def back_propagation(hidden_layer, output_layer, one_hot_encoding, learning_rate, x):
        output_derivative = np.zeros(2)
        output_gradient = np.zeros(2)
        for i in range(0, len(output_layer)):
            output_derivative[i] = (1.0 - output_layer[i]) * output_layer[i]
        for i in range(0, len(output_layer)):
            output_gradient[i] = output_derivative[i] * (one_hot_encoding[i] - output_layer[i])
        hidden_derivative = np.zeros(72)
        hidden_gradient = np.zeros(72)
        for i in range(0, len(hidden_layer)):
            hidden_derivative[i] = (1.0 - hidden_layer[i]) * (1.0 + hidden_layer[i])
        for i in range(0, len(hidden_layer)):
            sum_ = 0
            for j in range(0, len(output_gradient)):
                sum_ += output_gradient[j] * hidden_weights[i][j]
            hidden_gradient[i] = sum_ * hidden_derivative[i]
        recalculate_weights(learning_rate, hidden_weights, output_gradient, hidden_layer)
        recalculate_weights(learning_rate, weights, hidden_gradient, x)
```

Output

```
In [9]: one_hot_encoding = np.zeros((2,2))
for i in range(0, len(one_hot_encoding)):
    one_hot_encoding[i][i] = 1
training_correct_answers = 0
for i in range(0, len(x_train)):
    activate_layer(hidden_layer, weights, x_train[i])
    activate_layer(output_layer, hidden_weights, hidden_layer)
    output_layer = soft_max(output_layer)
    training_correct_answers += 1 if y_train[i] == np.argmax(output_layer) else 0
    back_propagation(hidden_layer, output_layer, one_hot_encoding[int(y_train[i])], -1, x_train[i])
print("MLP Correct answers while learning: %s / %s (Accuracy = %s) on %s database." % (training_correct_answers, len(x_train),
                                                                                       training_correct_answers/len(x_train), "Duke breast cancer database"))
```

MLP Correct answers while learning: 51 / 77 (Accuracy = 0.6623376623376623) on Duke breast cancer database.

```
In [10]: testing_correct_answers = 0
for i in range(0, len(x_test)):
    activate_layer(hidden_layer, weights, x_test[i])
    activate_layer(output_layer, hidden_weights, hidden_layer)
    output_layer = soft_max(output_layer)
    testing_correct_answers += 1 if y_test[i] == np.argmax(output_layer) else 0
print("MLP Correct answers while testing: %s / %s (Accuracy = %s) on %s database" % (testing_correct_answers, len(x_test),
                                                                                       testing_correct_answers/len(x_test), "Duke breast cancer database"))
```

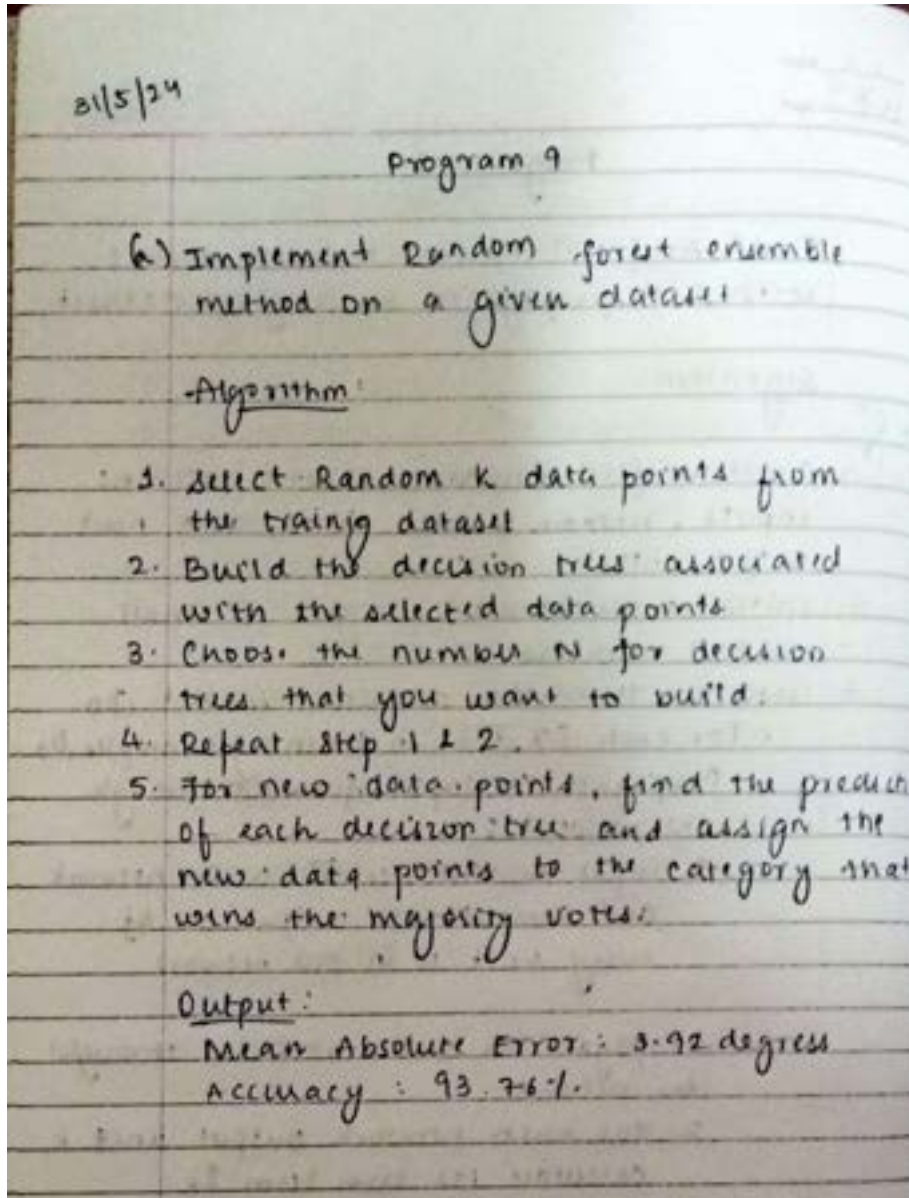
MLP Correct answers while testing: 9 / 9 (Accuracy = 1.0) on Duke breast cancer database

PROGRAM 9

Date: 31-05-2024

a) Implement Random forest ensemble method on a given dataset.

Algorithm



Code

[Open in Colab](#)

```

3> [1]: # Pandas is used for data manipulation
import pandas as pd
# Read in data and display first 5 rows
features = pd.read_csv('/content/temps.csv')
features.head()

```

```

Out[1]:
   year  month  day  week  temp_2  temp_1  average  actual  forecast_moa  forecast_act  forecast_under  friend
0  2016     1    1    Fri    45     45    45.5     45         43         50         44         29
1  2016     1    2    Sat    44     43    43.7     44         41         50         44         31
2  2016     1    3    Sun    45     44    45.8     41         43         48         47         38
3  2016     1    4    Mon    44     41    45.9     40         44         48         48         33
4  2016     1    5    Tue    41     40    45.0     44         48         48         48         41

```

```

3> [2]: print("The shape of our features is:", features.shape)

```

The shape of our features is: (348, 12)

```

3> [3]: # Descriptive statistics for each column
features.describe()

```

```

Out[3]:
   year  month  day  temp_2  temp_1  average  actual  forecast_moa  forecast_act  forecast_under  friend
count  348.0  348.000000  348.000000  348.000000  348.000000  348.000000  348.000000  348.000000  348.000000  348.000000  348.000000
mean   2016.0  6.477011  15.514268  52.862299  52.791148  59.760632  52.543153  57.236508  52.373940  58.772989  60.000000
std      0.0  8.438380  8.772952  12.168398  12.105642  10.527308  11.794148  10.805749  10.348881  10.705298  15.000000
min   2016.0  1.000000  1.000000  35.000000  38.000000  45.100000  39.000000  41.000000  48.000000  44.000000  28.000000
25%   2016.0  3.000000  8.000000  54.000000  54.000000  49.975000  54.000000  48.000000  53.000000  50.000000  47.000000
50%   2016.0  6.000000  15.000000  52.500000  52.500000  58.200000  52.500000  56.000000  51.000000  58.000000  60.000000
75%   2016.0  10.000000  23.000000  71.000000  71.000000  69.025000  71.000000  68.000000  72.000000  69.000000  71.000000
max   2016.0  12.000000  31.000000  117.000000  117.000000  77.400000  92.000000  77.000000  82.000000  79.000000  96.000000

```

◀ ▶

```

3> [4]: # One-hot encode the data using pandas get_dummies
features = pd.get_dummies(features)
# Display the first 5 rows of the last 32 columns
features.iloc[:,31:].head()

```

```

Out[4]:
   average  actual  forecast_moa  forecast_act  forecast_under  friend  week_Fri  week_Mon  week_Sat  week_Sun  week_Thurs  week_Tu
0    45.5     45         43         50         44         29         True         False         False         False         False         False
1    43.7     44         41         50         44         31         False         False         True         False         False         False
2    45.8     41         43         48         47         38         False         False         False         True         False         False
3    45.9     40         44         48         48         33         False         True         False         False         False         False
4    45.0     44         48         48         48         41         False         False         False         False         False         False

```

◀ ▶

```

In [5]: # Use numpy to convert to arrays
import numpy as np
# labels are the values we want to predict
labels = np.array(features['actual'])
# Remove the labels from the features
# axis 1 refers to the columns
features = features.drop('actual', axis = 1)
# Saving feature names for later use
feature_list = list(features.columns)
# Convert to numpy array
features = np.array(features)

In [6]: # Using Skikit-learn to split data into training and testing sets
from sklearn.model_selection import train_test_split
# Split the data into training and testing sets
train_features, test_features, train_labels, test_labels = train_test_split(features, labels, test_size = 0.25, random_state=42)

In [7]: print('Training Features Shape:', train_features.shape)
print('Training Labels Shape:', train_labels.shape)
print('Testing Features Shape:', test_features.shape)
print('Testing Labels Shape:', test_labels.shape)

Training Features Shape: (261, 17)
Training Labels Shape: (261,)
Testing Features Shape: (87, 17)
Testing Labels Shape: (87,)

In [8]: # The baseline predictions are the historical averages
baseline_preds = test_features[:, feature_list.index('average')]
# Baseline errors, and display average baseline error
baseline_errors = abs(baseline_preds - test_labels)
print('Average baseline error: ', round(np.mean(baseline_errors), 2))

Average baseline error: 5.05

In [9]: # Import the model we are using
from sklearn.ensemble import RandomForestRegressor
# Instantiate model with 1000 decision trees
rf = RandomForestRegressor(n_estimators = 1000, random_state = 42)
# Train the model on training data
rf.fit(train_features, train_labels);

In [10]: # Use the forest's predict method on the test data
predictions = rf.predict(test_features)
# Calculate the absolute errors
errors = abs(predictions - test_labels)
# Print out the mean absolute error (mae)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')

Mean Absolute Error: 3.87 degrees.

In [11]: # Calculate mean absolute percentage error (MAPE)
mape = 100 * (errors / test_labels)
# Calculate and display accuracy
accuracy = 100 - np.mean(mape)
print('Accuracy:', round(accuracy, 2), '%.')

Accuracy: 93.93 %.

```

```
In [12]: # Import tools needed for visualization
from sklearn.tree import export_graphviz
import pydot
# Pull out one tree from the forest
tree = rf.estimators_[5]
# Import tools needed for visualization
from sklearn.tree import export_graphviz
import pydot
# Pull out one tree from the forest
tree = rf.estimators_[5]
# Export the image to a dot file
export_graphviz(tree, out_file = 'tree.dot', feature_names = feature_list, rounded = True, precision = 1)
# Use dot file to create a graph
(graph, ) = pydot.graph_from_dot_file('tree.dot')
# Write graph to a png file
graph.write_png('tree.png')
```


```
In [13]: # Limit depth of tree to 3 levels
rf_small = RandomForestRegressor(n_estimators=10, max_depth = 3)
rf_small.fit(train_features, train_labels)
# Extract the small tree
tree_small = rf_small.estimators_[5]
# Save the tree as a png image
export_graphviz(tree_small, out_file = 'small_tree.dot', feature_names = feature_list, rounded = True, precision = 1)
(graph, ) = pydot.graph_from_dot_file('small_tree.dot')
graph.write_png('small_tree.png');
```

```
In [14]: # Get numerical feature importances
importances = list(rf.feature_importances_)
# List of tuples with variable and importance
feature_importances = [(feature, round(importance, 2)) for feature, importance in zip(feature_list, importances)]
# Sort the feature importances by most important first
feature_importances = sorted(feature_importances, key = lambda x: x[1], reverse = True)
# Print out the feature and importances
[print('Variable: {:20} Importance: {}'.format(pair[0], pair[1])) for pair in feature_importances];
```

Variable: temp_1	Importance: 0.66
Variable: average	Importance: 0.15
Variable: forecast_noaa	Importance: 0.05
Variable: forecast_acc	Importance: 0.03
Variable: day	Importance: 0.02
Variable: temp_2	Importance: 0.02
Variable: forecast_under	Importance: 0.02
Variable: friend	Importance: 0.02
Variable: month	Importance: 0.01
Variable: year	Importance: 0.0
Variable: week_Fri	Importance: 0.0
Variable: week_Mon	Importance: 0.0
Variable: week_Sat	Importance: 0.0
Variable: week_Sun	Importance: 0.0
Variable: week_Thurs	Importance: 0.0
Variable: week_Tues	Importance: 0.0
Variable: week_Wed	Importance: 0.0

```
In [15]: # New random forest with only the two most important variables
rf_most_important = RandomForestRegressor(n_estimators= 1000, random_state=42)
# Extract the two most important features
important_indices = [feature_list.index('temp_1'), feature_list.index('average')]
train_important = train_features[:, important_indices]
test_important = test_features[:, important_indices]
# Train the random forest
rf_most_important.fit(train_important, train_labels)
# Make predictions and determine the error
predictions = rf_most_important.predict(test_important)
errors = abs(predictions - test_labels)
# Display the performance metrics
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
mape = np.mean(100 * (errors / test_labels))
accuracy = 100 - mape
print('Accuracy:', round(accuracy, 2), '%.')
```


Output



b) Implement Boosting ensemble method on a given dataset.

Algorithm

(b) Implement Boosting Ensemble on a given dataset

Algorithm:

1. Initialize the dataset and assign equal weight to each of the data point.
2. Provide this as input to the model and identify the wrongly classified datapoint.
3. Increase the weight of the wrongly classified data points and decrease the weights of correctly classified data points. And then normalize the weights of all data points.
4. If (got required results)
 Goto step - 5
Else
 Goto step - 2
5. End.

Output

Confusion Matrix : $\begin{bmatrix} 116 & 35 \\ 26 & 54 \end{bmatrix}$

Accuracy Score : 0.7359.

Code

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
sns.set_style("whitegrid")
plt.style.use("fivethirtyeight")
```

```
In [2]: df = pd.read_csv("../content/diabetes.csv")
df.head()
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	140	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   Pregnancies           768 non-null   int64  
 1   Glucose               768 non-null   int64  
 2   BloodPressure         768 non-null   int64  
 3   SkinThickness        768 non-null   int64  
 4   Insulin              768 non-null   int64  
 5   BMI                  768 non-null   float64 
 6   DiabetesPedigreeFunction 768 non-null   float64 
 7   Age                  768 non-null   int64  
 8   Outcome              768 non-null   int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [4]: df.isnull().sum()
```

```
Out[4]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness 0
Insulin      0
BMI          0
DiabetesPedigreeFunction 0
Age          0
Outcome      0
dtype: int64
```

```
In [5]: pd.set_option('display.float_format', '{:.2f}'.format)
df.describe()
```

```
Out[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.00	768.00	768.00	768.00	768.00	768.00	768.00	768.00	768.00
mean	3.85	120.89	69.11	20.54	79.80	31.99	0.47	33.24	0.35
std	3.37	31.97	19.36	15.95	115.24	7.88	0.33	11.76	0.48
min	0.00	0.00	0.00	0.00	0.00	0.00	0.08	21.00	0.00
25%	1.00	99.00	62.00	0.00	0.00	27.30	0.24	24.00	0.00
50%	3.00	117.00	72.00	23.00	30.90	32.00	0.37	29.00	0.00
75%	6.00	140.25	80.00	32.00	127.25	36.60	0.63	41.00	1.00
max	17.00	199.00	122.00	99.00	846.00	67.10	2.42	81.00	1.00

```

In [6]: categorical_val = []
        continuous_val = []
        for column in df.columns:
            # print("=====")
            # print(f"{column} : {df[column].unique()}")
            if len(df[column].unique()) <= 10:
                categorical_val.append(column)
            else:
                continuous_val.append(column)

In [7]: df.columns

Out[7]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
              'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')

In [8]: # How many missing zeros are missing in each feature
        feature_columns = [
            'Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
            'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age'
        ]

        for column in feature_columns:
            print("=====")
            print(f"{column} ==> Missing zeros : {len(df.loc[df[column] == 0])}")

=====
Pregnancies ==> Missing zeros : 111
=====
Glucose ==> Missing zeros : 5
=====
BloodPressure ==> Missing zeros : 35
=====
SkinThickness ==> Missing zeros : 227
=====
Insulin ==> Missing zeros : 374
=====
BMI ==> Missing zeros : 11
=====
DiabetesPedigreeFunction ==> Missing zeros : 0
=====
Age ==> Missing zeros : 0

In [9]: from sklearn.impute import SimpleImputer

        fill_values = SimpleImputer(missing_values=0, strategy="mean", copy=False)
        df[feature_columns] = fill_values.fit_transform(df[feature_columns])

        for column in feature_columns:
            print("=====")
            print(f"{column} ==> Missing zeros : {len(df.loc[df[column] == 0])}")

=====
Pregnancies ==> Missing zeros : 0
=====
Glucose ==> Missing zeros : 0
=====
BloodPressure ==> Missing zeros : 0
=====
SkinThickness ==> Missing zeros : 0
=====
Insulin ==> Missing zeros : 0
=====
BMI ==> Missing zeros : 0
=====
DiabetesPedigreeFunction ==> Missing zeros : 0
=====
Age ==> Missing zeros : 0

```

```
In [10]: from sklearn.model_selection import train_test_split

X = df[feature_columns]
y = df.Outcome

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

In [11]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_report

def evaluate(model, X_train, X_test, y_train, y_test):
    y_test_pred = model.predict(X_test)
    y_train_pred = model.predict(X_train)

    print("TRAINING RESULTS: \n\n")
    clf_report = pd.DataFrame(classification_report(y_train, y_train_pred, output_dict=True))
    print(f"CONFUSION MATRIX:\n{confusion_matrix(y_train, y_train_pred)}")
    print(f"ACCURACY SCORE:\n{accuracy_score(y_train, y_train_pred):.4f}")
    print(f"CLASSIFICATION REPORT:\n{clf_report}")

    print("TESTING RESULTS: \n\n")
    clf_report = pd.DataFrame(classification_report(y_test, y_test_pred, output_dict=True))
    print(f"CONFUSION MATRIX:\n{confusion_matrix(y_test, y_test_pred)}")
    print(f"ACCURACY SCORE:\n{accuracy_score(y_test, y_test_pred):.4f}")
    print(f"CLASSIFICATION REPORT:\n{clf_report}")

In [12]: from sklearn.ensemble import AdaBoostClassifier

ada_boost_clf = AdaBoostClassifier(n_estimators=30)
ada_boost_clf.fit(X_train, y_train)
evaluate(ada_boost_clf, X_train, X_test, y_train, y_test)
```

Output-AdaBoost

```
TRAINING RESULTS:
=====
CONFUSION MATRIX:
[[310  39]
 [ 51 137]]
ACCURACY SCORE:
0.8324
CLASSIFICATION REPORT:

```

	0	1	accuracy	macro avg	weighted avg
precision	0.86	0.78	0.83	0.82	0.83
recall	0.89	0.73	0.83	0.81	0.83
f1-score	0.87	0.75	0.83	0.81	0.83
support	349.00	188.00	0.83	537.00	537.00

```
TESTING RESULTS:
=====
CONFUSION MATRIX:
[[123  28]
 [ 27  53]]
ACCURACY SCORE:
0.7619
CLASSIFICATION REPORT:

```

	0	1	accuracy	macro avg	weighted avg
precision	0.82	0.65	0.76	0.74	0.76
recall	0.81	0.66	0.76	0.74	0.76
f1-score	0.82	0.66	0.75	0.74	0.76
support	151.00	80.00	0.76	231.00	231.00

Output- GradientBoost

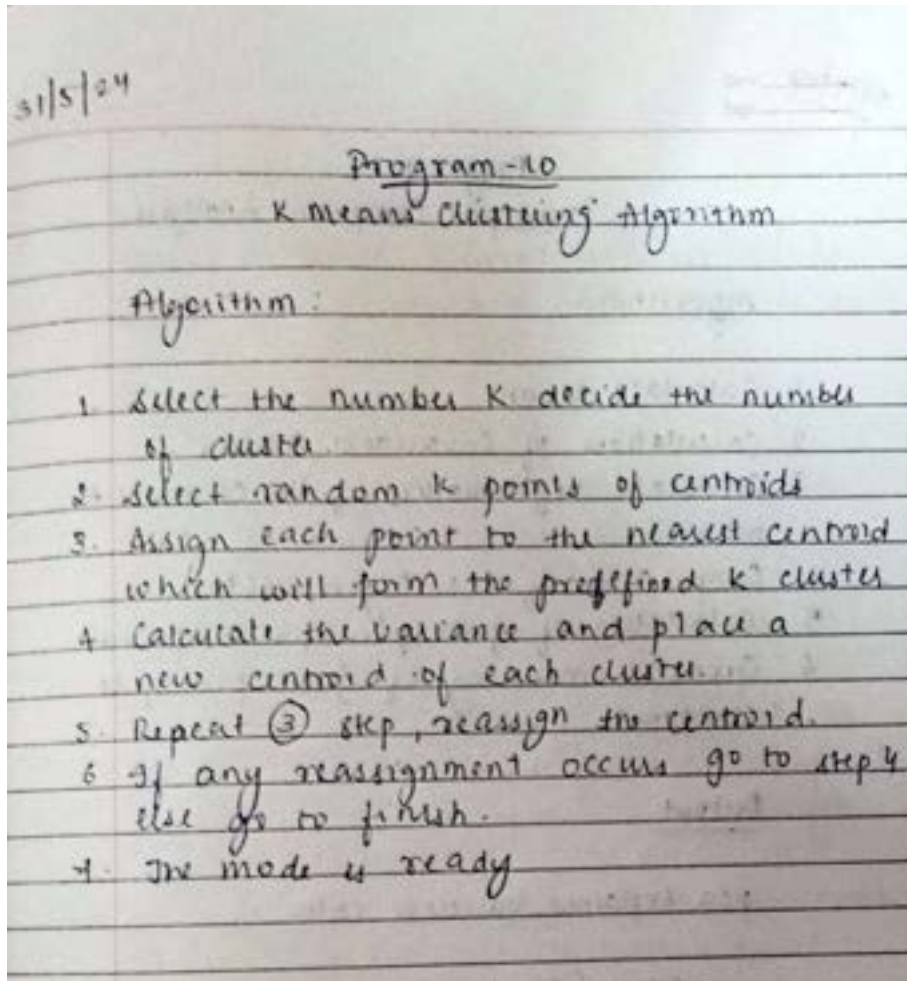
```
TRAINING RESULTS:
*****
CONFUSION MATRIX:
[[342  7]
 [ 19 169]]
ACCURACY SCORE:
0.9516
CLASSIFICATION REPORT:
      _
precision  0.95  0.96      0.95  macro avg  weighted avg
recall    0.98  0.90      0.95  0.94      0.95
f1-score   0.96  0.93      0.95  0.95      0.95
support   349.00 188.00      0.95  537.00      537.00
TESTING RESULTS:
*****
CONFUSION MATRIX:
[[116  35]
 [ 26  54]]
ACCURACY SCORE:
0.7359
CLASSIFICATION REPORT:
      _
precision  0.82  0.61      0.74  0.71      0.74
recall    0.77  0.68      0.74  0.72      0.74
f1-score   0.79  0.64      0.74  0.72      0.74
support   151.00  80.00      0.74  231.00      231.00
```

PROGRAM 10

Date: 24-05-2024

Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Algorithm



Code

Importing and initializing the data points

```
In [5]: import matplotlib.pyplot as plt
import numpy as np
from sklearn.cluster import KMeans

In [6]: from sklearn.datasets import make_blobs
X, y_true = make_blobs(n_samples=500, centers=4, cluster_std=0.60, random_state=0)

In [7]: import plotly.express as px
fig = px.scatter(x=X[:, 0], y=X[:, 1], width=800, height=500)
fig.show()
```

Elbow Method to find optimal K

```
In [8]: cost = []
for i in range(1, 11):
    KM = KMeans(n_clusters = i, max_iter = 500)
    KM.fit(X)
    cost.append(KM.inertia_)


# plot the cost against K values
fig = px.line(x=range(1, 11), y=cost, width=800, height=400)
fig.show()
# the point of the elbow is the
# most optimal value for choosing K
```

Defining Model and fitting the same

```
In [9]: kmeans = KMeans(n_clusters=4)
kmeans.fit(X)
y_kmeans = kmeans.predict(X)

In [10]: fig = px.scatter(x=X[:, 0], y=X[:, 1], color=y_kmeans, width=700, height=400)
trace = px.scatter(x=X[:, 0], y=X[:, 1], width=700, height=400)
fig.show()

In [11]: plt.scatter(X[:, 0], X[:, 1], c=y_kmeans, s=20)
centers = kmeans.cluster_centers_
plt.scatter(centers[:, 0], centers[:, 1], c='blue', s=100, alpha=0.5);
plt.show()
```



Iris Dataset

```
In [14]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import datasets

In [15]: iris = datasets.load_iris()
df = pd.DataFrame(iris.data)
df['class'] = iris.target
df.columns = ['sepal_len', 'sepal_wid', 'petal_len', 'petal_wid', 'class']
df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   column      Non-Null count  Dtype
---  --
 0   sepal_len    150 non-null         float64
 1   sepal_wid    150 non-null         float64
 2   petal_len    150 non-null         float64
 3   petal_wid    150 non-null         float64
 4   class        150 non-null         int64
dtypes: float64(4), int64(1)
memory usage: 8.9 KB

In [16]: px.histogram(df, x = 'class', color='class')

In [17]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
x = df.iloc[:,0:4].values

In [18]: scaled_x = scaler.fit_transform(x)

In [19]: model = KMeans(n_clusters=3, init='k-means++', random_state=0)
labels = model.fit_predict(scaled_x)

In [20]: import matplotlib.pyplot as plt
fig = plt.figure()

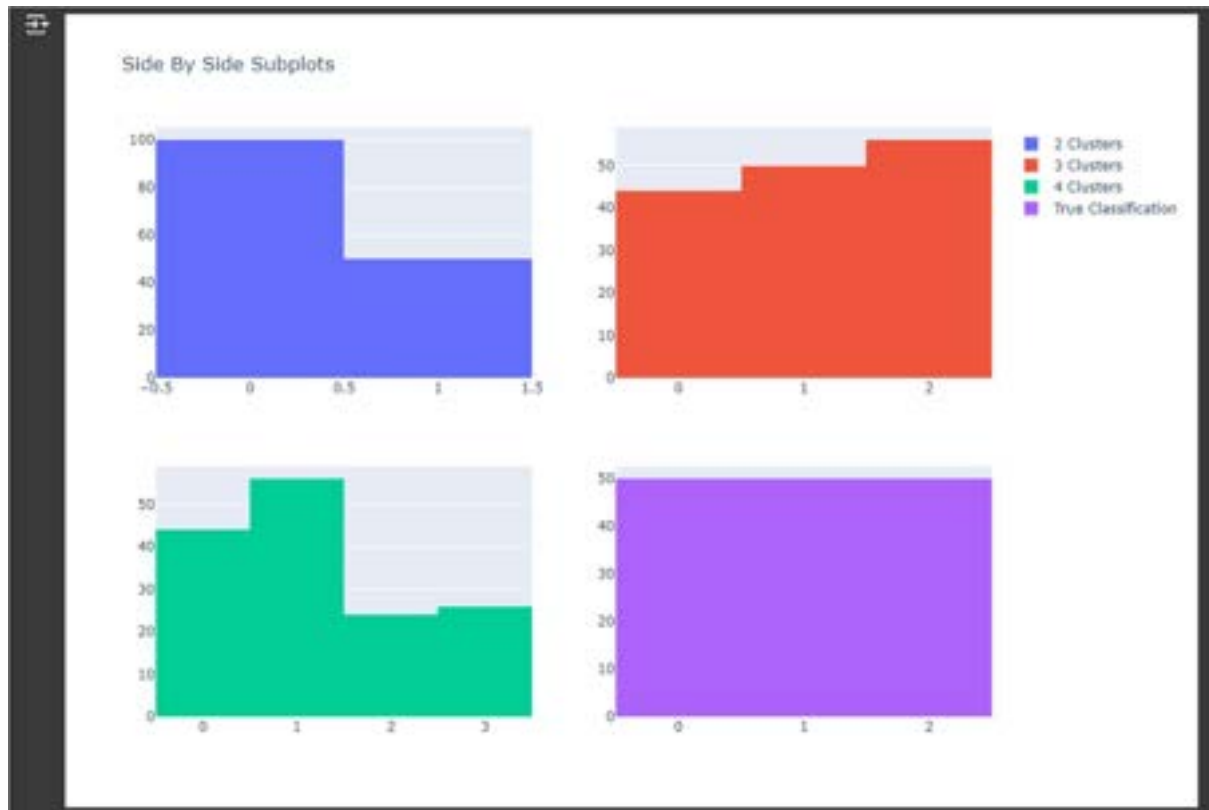
# Add trace
fig.add_trace(px.histogram(x=labels, name="Predicted Labels"))
fig.add_trace(px.histogram(x=df['class'], name="True Labels"))

# Overlay both histograms
fig.update_layout(barnode="overlay")
# Reduce opacity to see both histograms
fig.update_traces(opacity=0.75)
fig.show()

In [21]: labels = []
for i in range(2, 5):
    model = KMeans(n_clusters = 1, max_iter = 500)
    model.fit(scaled_x)
    labels.append(model.fit_predict(scaled_x))

In [22]: from matplotlib.figure import Figure
import matplotlib.pyplot as plt
fig = Figure(figsize=(10, 10))
for i in range(3, 5):
    fig.add_subplot(2, 2, i+1, title="Cluster {}".format(i+1))
    fig.add_subplot(2, 2, i+2, title="True Classification")
    fig.add_subplot(2, 2, i+3, title="True Classification")
    fig.add_subplot(2, 2, i+4, title="True Classification")
    fig.show()
```

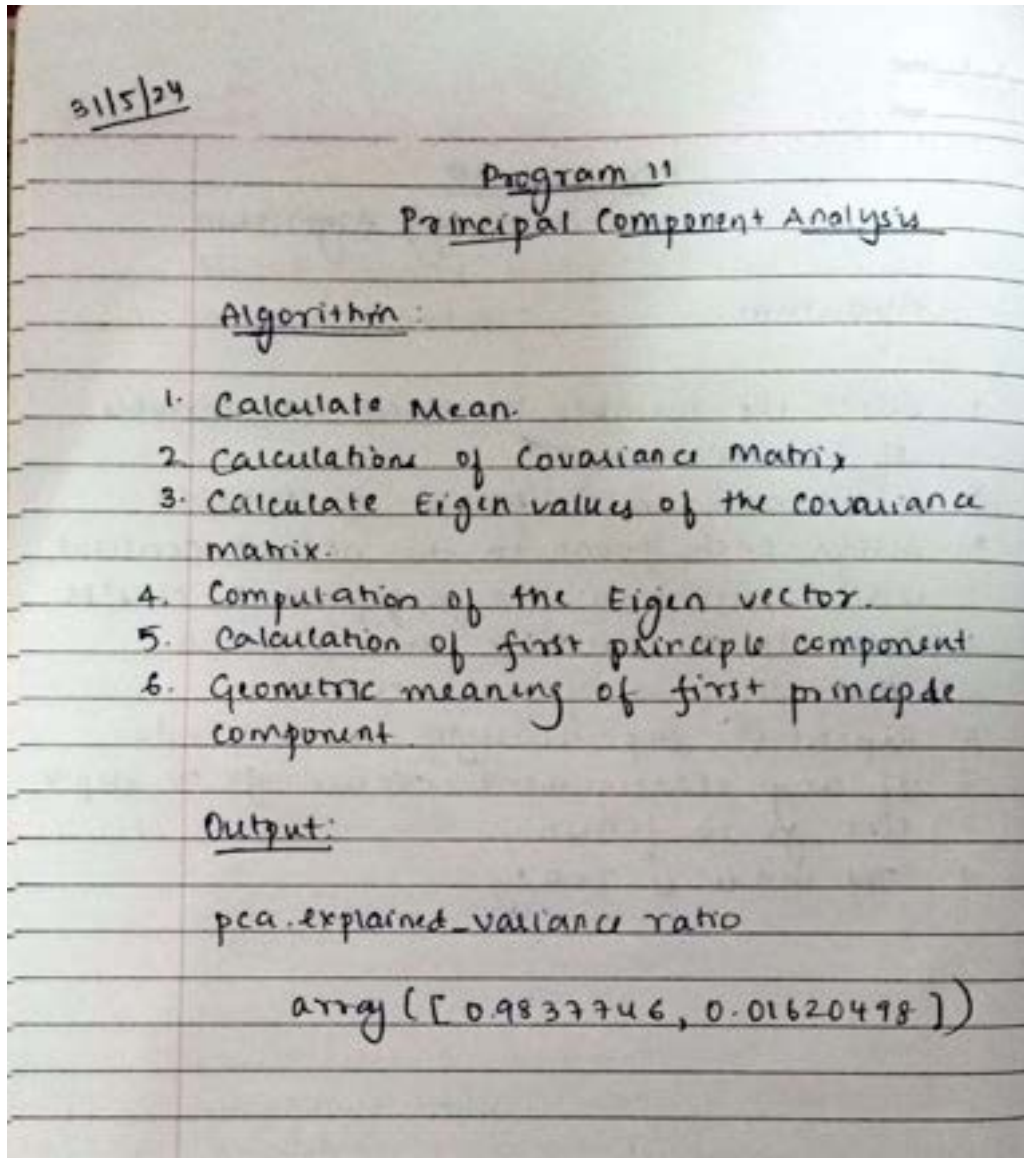
Output



PROGRAM 11

Date: 24-05-2024

Implement Dimensionality reduction using Principle Component Analysis (PCA) method.



Code

Open in Colab

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

In [ ]: import seaborn as sns
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots

In [ ]: df = pd.read_csv('/content/drive/MyDrive/breast-cancer.csv')
df.head()

Out[ ]:
   id diagnosis radius_mean texture_mean perimeter_mean area_mean smoothness_mean compactness_mean concavity_mean
0  842302      M      17.99       10.38        122.80      1001.0         0.11840         0.27760         0.3001
1  842517      M      20.57       17.77        132.90      1326.0         0.08474         0.07864         0.0866
2  84300908     M      19.69       21.25        130.00      1203.0         0.10960         0.18990         0.1974
3  84348301      M      11.42       20.38         77.58        386.1         0.14250         0.28390         0.2414
4  84358402      M      20.29       14.34        135.10      1297.0         0.10030         0.13280         0.1980

5 rows x 10 columns

4

In [ ]: df.drop('id', axis=1, inplace=True) #drop redundant columns

In [ ]: df['diagnosis'] = (df['diagnosis'] == 'M').astype(int) #encode the label into 1/0

In [ ]: corr = df.corr()

In [ ]: # Get the absolute value of the correlation
cor_target = abs(corr['diagnosis'])

# Select highly correlated features (threshold = 0.2)
relevant_features = cor_target[cor_target>0.2]

# Collect the names of the features
names = [index for index, value in relevant_features.items()]

# Drop the target variable from the results
names.remove('diagnosis')

# Display the results
print(names)

['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean', 'concave points_mean', 'symmetry_mean', 'radius_se', 'perimeter_se', 'area_se', 'compactness_se', 'concavity_se', 'concave points_se', 'radius_worst', 'texture_worst', 'perimeter_worst', 'area_worst', 'smoothness_worst', 'compactness_worst', 'concavity_worst', 'concave points_worst', 'symmetry_worst', 'fractal_dimension_worst']

In [ ]: X = df[names].values
```

```

In [ ]: class PCA:
    """
    Principal Component Analysis (PCA) class for dimensionality reduction.
    """

    def __init__(self, n_components):
        """
        Constructor method that initializes the PCA object with the number of components to retain.

        Args:
        - n_components (int): Number of principal components to retain.
        """
        self.n_components = n_components

    def fit(self, X):
        """
        Fits the PCA model to the input data and computes the principal components.

        Args:
        - X (numpy.ndarray): Input data matrix with shape (n_samples, n_features).
        """
        # Compute the mean of the input data along each feature dimension.
        mean = np.mean(X, axis=0)

        # Subtract the mean from the input data to center it around zero.
        X = X - mean

        # Compute the covariance matrix of the centered input data.
        cov = np.cov(X.T)

        # Compute the eigenvectors and eigenvalues of the covariance matrix.
        eigenvalues, eigenvectors = np.linalg.eigh(cov)
        # Reverse the order of the eigenvalues and eigenvectors.
        eigenvalues = eigenvalues[::-1]
        eigenvectors = eigenvectors[:, ::-1]

        # Keep only the first n_components eigenvectors as the principal components.
        self.components = eigenvectors[:, :self.n_components]

        # Compute the explained variance ratio for each principal component.
        # Compute the total variance of the input data
        total_variance = np.sum(np.var(X, axis=0))

        # Compute the variance explained by each principal component
        self.explained_variances = eigenvalues[:self.n_components]

        # Compute the explained variance ratio for each principal component
        self.explained_variance_ratio_ = self.explained_variances / total_variance

    def transform(self, X):
        """
        Transforms the input data by projecting it onto the principal components.

        Args:
        - X (numpy.ndarray): Input data matrix with shape (n_samples, n_features).

        Returns:
        - transformed_data (numpy.ndarray): Transformed data matrix with shape (n_samples, n_components).
        """
        # Center the input data around zero using the mean computed during the fit step.
        X = X - np.mean(X, axis=0)

        # Project the centered input data onto the principal components.
        transformed_data = np.dot(X, self.components)

        return transformed_data

    def fit_transform(self, X):
        """
        Fits the PCA model to the input data and computes the principal components then
        transforms the input data by projecting it onto the principal components.

        Args:
        - X (numpy.ndarray): Input data matrix with shape (n_samples, n_features).
        """
        self.fit(X)
        transformed_data = self.transform(X)
        return transformed_data

```

```

In [ ]: pca = PCA(2)

In [ ]: pca.fit(X)

In [ ]: pca.explained_variance_ratio_

Out[ ]: array([0.98377428, 0.01620498])

In [ ]: X_transformed = pca.transform(X)

In [ ]: X_transformed[:,1].shape

Out[ ]: (569,)

In [ ]: fig = plt.scatter(x=X_transformed[:,0], y=X_transformed[:,1])
fig.update_layout(
    title="PCA transformed data for breast cancer dataset",
    xaxis_title="PC1",
    yaxis_title="PC2"
)
fig.show()

```

Output

